



```
import sys
in_colab = 'google.colab' in sys.modules
if in_colab:
    # Install packages in Colab
    !pip install category_encoders==2.0.0
    !pip install pandas-profiling==2.3.0
    !pip install plotly==4.1.1
```



```

Requirement already satisfied: category_encoders==2.0.0 in /usr/local/lib/python3.6/dist-packages (2.0.0)
Requirement already satisfied: scipy>=0.19.0 in /usr/local/lib/python3.6/dist-packages (from category_encoders=
Requirement already satisfied: patsy>=0.4.1 in /usr/local/lib/python3.6/dist-packages (from category_encoders==
Requirement already satisfied: numpy>=1.11.3 in /usr/local/lib/python3.6/dist-packages (from category_encoders=
Requirement already satisfied: statsmodels>=0.6.1 in /usr/local/lib/python3.6/dist-packages (from category_encoc
Requirement already satisfied: pandas>=0.21.1 in /usr/local/lib/python3.6/dist-packages (from category_encoders
Requirement already satisfied: scikit-learn>=0.20.0 in /usr/local/lib/python3.6/dist-packages (from category_er
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from patsy>=0.4.1->category_encoc
Requirement already satisfied: pytz>=2011k in /usr/local/lib/python3.6/dist-packages (from pandas>=0.21.1->cate
Requirement already satisfied: python-dateutil>=2.5.0 in /usr/local/lib/python3.6/dist-packages (from pandas>=0.
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.6/dist-packages (from scikit-learn>=0.20.
Collecting pandas-profiling==2.3.0
Requirement already satisfied: astropy in /usr/local/lib/python3.6/dist-packages (from pandas-profiling==2.3.0)
Requirement already satisfied: missingno>=0.4.2 in /usr/local/lib/python3.6/dist-packages (from pandas-profilir
Collecting confuse>=1.0.0 (from pandas-profiling==2.3.0)
Requirement already satisfied: jinja2>=2.8 in /usr/local/lib/python3.6/dist-packages (from pandas-profiling==2.
Collecting phik>=0.9.8 (from pandas-profiling==2.3.0)
Using cached https://files.pythonhosted.org/packages/45/ad/24a16fa4ba612fb96a3c4bb115a5b9741483f53b66d3d3afd5
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.6/dist-packages (from pandas-profiling==2
Requirement already satisfied: matplotlib>=1.4 in /usr/local/lib/python3.6/dist-packages (from pandas-profiling
Requirement already satisfied: htmlmin>=0.1.12 in /usr/local/lib/python3.6/dist-packages (from pandas-profiling
Requirement already satisfied: numpy>=1.10.0 in /usr/local/lib/python3.6/dist-packages (from astropy->pandas-pr
Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from missingno>=0.4.2->pandas-p
Requirement already satisfied: seaborn in /usr/local/lib/python3.6/dist-packages (from missingno>=0.4.2->pandas
Requirement already satisfied: pyyaml in /usr/local/lib/python3.6/dist-packages (from confuse>=1.0.0->pandas-pr
Requirement already satisfied: MarkupSafe>=0.23 in /usr/local/lib/python3.6/dist-packages (from jinja2>=2.8->pa
Requirement already satisfied: nbconvert>=5.3.1 in /usr/local/lib/python3.6/dist-packages (from phik>=0.9.8->pa
Collecting pytest-pylint>=0.13.0 (from phik>=0.9.8->pandas-profiling==2.3.0)
Using cached https://files.pythonhosted.org/packages/64/dc/6f35f114844fb12e38d60c4f3d2441a55baff7043ad4e01377
Requirement already satisfied: jupyter-client>=5.2.3 in /usr/local/lib/python3.6/dist-packages (from phik>=0.9.
Collecting pytest>=4.0.2 (from phik>=0.9.8->pandas-profiling==2.3.0)
Using cached https://files.pythonhosted.org/packages/0c/91/d68f68ce54cd3e8afa1ef73ea1ad44df2438521b64c0820e5f
Requirement already satisfied: numba>=0.38.1 in /usr/local/lib/python3.6/dist-packages (from phik>=0.9.8->panda
Requirement already satisfied: python-dateutil>=2.5.0 in /usr/local/lib/python3.6/dist-packages (from pandas>=0
Requirement already satisfied: pytz>=2011k in /usr/local/lib/python3.6/dist-packages (from pandas>=0.19->pandas
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib>=1.
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib>=1.4->pa
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packag
Requirement already satisfied: bleach in /usr/local/lib/python3.6/dist-packages (from nbconvert>=5.3.1->phik>=0
Requirement already satisfied: jupyter-core in /usr/local/lib/python3.6/dist-packages (from nbconvert>=5.3.1->p
Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.6/dist-packages (from nbconvert>=5.3.1-
Requirement already satisfied: pandocfilters>=1.4.1 in /usr/local/lib/python3.6/dist-packages (from nbconvert>=
Requirement already satisfied: testpath in /usr/local/lib/python3.6/dist-packages (from nbconvert>=5.3.1->phik>
Requirement already satisfied: pygments in /usr/local/lib/python3.6/dist-packages (from nbconvert>=5.3.1->phik>
Requirement already satisfied: defusedxml in /usr/local/lib/python3.6/dist-packages (from nbconvert>=5.3.1->phi
Requirement already satisfied: mistune<2,>=0.8.1 in /usr/local/lib/python3.6/dist-packages (from nbconvert>=5.3
Requirement already satisfied: nbformat>=4.4 in /usr/local/lib/python3.6/dist-packages (from nbconvert>=5.3.1->
Requirement already satisfied: entrypoints>=0.2.2 in /usr/local/lib/python3.6/dist-packages (from nbconvert>=5.
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from pytest-pylint>=0.13.0->phik>
Collecting pylint>=1.4.5 (from pytest-pylint>=0.13.0->phik>=0.9.8->pandas-profiling==2.3.0)
Using cached https://files.pythonhosted.org/packages/ea/f1/758de486e46ea2b8717992704b0fdd968b7c9b2bc790b976f6
Requirement already satisfied: pyzmq>=13 in /usr/local/lib/python3.6/dist-packages (from jupyter-client>=5.2.3
Requirement already satisfied: tornado>=4.1 in /usr/local/lib/python3.6/dist-packages (from jupyter-client>=5.2
Requirement already satisfied: py>=1.5.0 in /usr/local/lib/python3.6/dist-packages (from pytest>=4.0.2->phik>=0
Requirement already satisfied: importlib-metadata>=0.12; python_version < "3.8" in /usr/local/lib/python3.6/dis
Requirement already satisfied: packaging in /usr/local/lib/python3.6/dist-packages (from pytest>=4.0.2->phik>=0
Requirement already satisfied: attrs>=17.4.0 in /usr/local/lib/python3.6/dist-packages (from pytest>=4.0.2->phi
Requirement already satisfied: pluggy<1.0,>=0.12 in /usr/local/lib/python3.6/dist-packages (from pytest>=4.0.2-
Requirement already satisfied: more-itertools>=4.0.0 in /usr/local/lib/python3.6/dist-packages (from pytest>=4.
Requirement already satisfied: wcwidth in /usr/local/lib/python3.6/dist-packages (from pytest>=4.0.2->phik>=0.9
Requirement already satisfied: atomicwrites>=1.0 in /usr/local/lib/python3.6/dist-packages (from pytest>=4.0.2-
Requirement already satisfied: llvmlite>=0.25.0dev0 in /usr/local/lib/python3.6/dist-packages (from numba>=0.38
Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from kiwisolver>=1.0.1->ma
Requirement already satisfied: webencodings in /usr/local/lib/python3.6/dist-packages (from bleach->nbconvert>=
Requirement already satisfied: decorator in /usr/local/lib/python3.6/dist-packages (from traitlets>=4.2->nbconv
Requirement already satisfied: ipython-genutils in /usr/local/lib/python3.6/dist-packages (from traitlets>=4.2-
Requirement already satisfied: jsonschema!=2.5.0,>=2.4 in /usr/local/lib/python3.6/dist-packages (from nbformat
Collecting mccabe<0.7,>=0.6 (from pylint>=1.4.5->pytest-pylint>=0.13.0->phik>=0.9.8->pandas-profiling==2.3.0)
Using cached https://files.pythonhosted.org/packages/87/89/479dc97e18549e21354893e4ee4ef36bd1d237534982482c3f
Collecting astroid<2.4,>=2.3.0 (from pylint>=1.4.5->pytest-pylint>=0.13.0->phik>=0.9.8->pandas-profiling==2.3.0)
Using cached https://files.pythonhosted.org/packages/64/d3/4ba68bd56297556c9c7a5072d71d1664feaa86d9726c737a9f

```

```


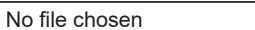
Using cached https://files.pythonhosted.org/packages/57/b7/45050082122070423f48407c1400007163ee770f/
Collecting isort<5,>=4.2.5 (from pylint>=1.4.5->pytest-pylint>=0.13.0->phik>=0.9.8->pandas-profiling==2.3.0)
Using cached https://files.pythonhosted.org/packages/e5/b0/c121fd1fa3419ea9bfd55c7f9c4fedfec5143208d8c7ad3ce3/
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.6/dist-packages (from importlib-metadata>=0.
Requirement already satisfied: wrapt==1.11.* in /usr/local/lib/python3.6/dist-packages (from astroid<2.4,>=2.3.
Collecting lazy-object-proxy==1.4.* (from astroid<2.4,>=2.3.0->pylint>=1.4.5->pytest-pylint>=0.13.0->phik>=0.9.
Using cached https://files.pythonhosted.org/packages/0e/26/534a6d32572a9dbca11619321535c0a7ab34688545d9d67c2c/
Collecting typed-ast<1.5,>=1.4.0; implementation_name == "cpython" and python_version < "3.8" (from astroid<2.4
Using cached https://files.pythonhosted.org/packages/31/d3/9d1802c161626d0278bafb1fffb32f76b9d01e123881bbf9d9/
ERROR: datascience 0.10.6 has requirement folium==0.2.1, but you'll have folium 0.8.3 which is incompatible.
Installing collected packages: confuse, pytest, mccabe, lazy-object-proxy, typed-ast, astroid, isort, pylint, p
Found existing installation: pytest 3.6.4
Uninstalling pytest-3.6.4:
Successfully uninstalled pytest-3.6.4
Found existing installation: pandas-profiling 1.4.1
Uninstalling pandas-profiling-1.4.1:
Successfully uninstalled pandas-profiling-1.4.1
Successfully installed astroid-2.3.2 confuse-1.0.0 isort-4.3.21 lazy-object-proxy-1.4.2 mccabe-0.6.1 pandas-prc
Requirement already satisfied: plotly==4.1.1 in /usr/local/lib/python3.6/dist-packages (4.1.1)
Requirement already satisfied: retrying>=1.3.3 in /usr/local/lib/python3.6/dist-packages (from plotly==4.1.1) (
Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from plotly==4.1.1) (1.12.0)

```

```

#Fetch smoking data file
from google.colab import files
uploaded = files.upload()


```



 Upload widget is only available when the cell has been executed in the current browser session.
 Saving cancerxx - for\_import.csv to cancerxx - for\_import.csv

```

# Load smoking data
import pandas as pd
import io
df_smoking = pd.read_csv(io.StringIO(uploaded['cancerxx - for_import.csv'].decode('utf-8')))
df_smoking.head()

```



	language	cereal_serve_per_month	cereal_times_per_month	more_than_one_cereal_type	milk_serve_per_month
0	5	3	2	2.0	3
1	4	0	0	NaN	0
2	5	5	2	2.0	5
3	3	1	1	2.0	4
4	5	2	2	1.0	0

 5 rows x 92 columns

```

# We assess the contents of df_smoking
df_smoking_shape = df_smoking.shape
print ('df_smoking Shape')
print (df_smoking_shape, '\n')
print ('df_smoking Count')
print (df_smoking.count(), '\n')
print ('df_smoking NaN Count')
print (df_smoking.isna().sum(), '\n')
print ('df_smoking Describe')
print (df_smoking.describe())

```



```
df_smoking Shape
(33672, 92)
```

```
df_smoking Count
language                33672
cereal_serve_per_month  33672
cereal_times_per_month  33672
more_than_one_cereal_type 22858
milk_serve_per_month    33672
milk_times_per_month    33672
milk_type               24044
soda_serve_per_month    33672
soda_times_per_month    33672
juice_serve_per_month   33672
juice_times_per_month   33672
coffee_serve_per_month  33672
coffee_times_per_month  33672
sports_drink_serve_per_month 33672
sports_drink_times_per_month 33672
fruit_drink_serve_per_month 33672
fruit_drink_times_per_month 33672
fruit_eat_serve_per_month 33672
fruit_eat_times_per_month 33672
salad_eat_serve_per_month 33672
salad_eat_times_per_month 33672
fries_eat_serve_per_month 33672
fries_eat_times_per_month 33672
potatoe_eat_serve_per_month 33672
potatoe_eat_times_per_month 33672
beans_eat_serve_per_month 33672
beans_eat_times_per_month 33672
grains_eat_serve_per_month 33672
grains_eat_times_per_month 33672
vegies_eat_serve_per_month 33672
...
vitD_reason            6906
1st_kind_cereal_eaten  22858
2nd_kind_cereal_eaten   9958
walk_past_wk           33672
walk_number_wk         10246
single_walk_distance   10229
single_walk_time       10229
walk_leisure_past_wk   32778
walk_leisure_number_wk 16074
walk_leisure_distance  16055
walk_leisure_time      16055
see_walking_from_home  33672
weather_discourages_walk 33672
walkway_existence      33672
walkable_retail        33672
walkable_bus_stop      33672
walkable_entertainment 33672
walkable_relaxation     33672
streets_have_walkways  33672
traffic_discourages_walking 33672
crime_discourages_walking 33672
animals_discourage_walking 33672
cigarette_even_once    33672
cigar_even_once        33672
pipe_even_once         33672
smokeless_even_once    33672
had_genetic_counseling 33672
genetic_counseling_with_MD 33672
genetic_counseling_for_cancer 33672
cigarettes_per_day      7602
Length: 92, dtype: int64
```

```
df_smoking NaN Count
language                0
cereal_serve_per_month  0
cereal times per month  0
```

```

cereal_times_per_month      0
more_than_one_cereal_type   10814
milk_serve_per_month        0
milk_times_per_month        0
milk_type                    9628
soda_serve_per_month        0
soda_times_per_month        0
juice_serve_per_month       0
juice_times_per_month       0
coffee_serve_per_month     0
coffee_times_per_month     0
sports_drink_serve_per_month 0
sports_drink_times_per_month 0
fruit_drink_serve_per_month 0
fruit_drink_times_per_month 0
fruit_eat_serve_per_month   0
fruit_eat_times_per_month   0
salad_eat_serve_per_month   0
salad_eat_times_per_month   0
fries_eat_serve_per_month   0
fries_eat_times_per_month   0
potatoe_eat_serve_per_month 0
potatoe_eat_times_per_month 0
beans_eat_serve_per_month   0
beans_eat_times_per_month   0
grains_eat_serve_per_month  0
grains_eat_times_per_month  0
vegies_eat_serve_per_month  0

```

...

```

vitD_reason      26766
1st_kind_cereal_eaten 10814
2nd_kind_cereal_eaten 23714
walk_past_wk      0
walk_number_wk    23426
single_walk_distance 23443
single_walk_time  23443
walk_leisure_past_wk 894
walk_leisure_number_wk 17598
walk_leisure_distance 17617
walk_leisure_time  17617
see_walking_from_home 0
weather_discourages_walk 0
walkway_existence 0
walkable_retail 0
walkable_bus_stop 0
walkable_entertainment 0
walkable_relaxation 0
streets_have_walkways 0
traffic_discourages_walking 0
crime_discourages_walking 0
animals_discourage_walking 0
cigarette_even_once 0
cigar_even_once 0
pipe_even_once 0
smokeless_even_once 0
had_genetic_counseling 0
genetic_counseling_with_MD 0
genetic_counseling_for_cancer 0
cigarettes_per_day 26070
Length: 92, dtype: int64

```

```

df_smoking Describe
      language ... cigarettes_per_day
count  33672.000000 ...      7602.000000
mean    4.670587 ...      22.540647
std     1.191156 ...      26.525465
min     1.000000 ...      1.000000
25%     4.000000 ...      6.000000
50%     5.000000 ...     15.000000
75%     5.000000 ...     20.000000
max     9.000000 ...     99.000000

```

[8 rows x 92 columns]

```
# Replace NaN to improve data format
import numpy as np
df_smoking1 = df_smoking.replace ({np.NaN: 0})
df_smoking1.head()
```

↗

	language	cereal_serve_per_month	cereal_times_per_month	more_than_one_cereal_type	milk_serve_per_month	
0	5	3	2	2.0	3	
1	4	0	0	0.0	0	
2	5	5	2	2.0	5	
3	3	1	1	2.0	4	
4	5	2	2	1.0	0	

5 rows × 92 columns

```
# Set up boolean columns such that yes = 1 and no = 0
features1 = {'more_than_one_cereal_type', 'vitamin_past_month', 'multivitamin_past_month', 'calcium_past_month', 'vitD_past_',
'walkway_existence', 'walkable_retail', 'walkable_bus_stop', 'walkable_entertainment', 'walkable_relaxation', 's',
'crime_discourages_walking', 'animals_discourage_walking', 'cigarette_even_once', 'cigar_even_once', 'pipe_even_',
'had_genetic_counseling', 'genetic_counseling_with_MD', 'genetic_counseling_for_cancer'}

replacements1 = {
    2: 0,
    3: 0,
    4: 0,
    5: 0,
    6: 0,
    7: 0,
    8: 0,
    9: 0
}

df_smoking2 = df_smoking1[features1].replace(replacements1)
df_smoking2.head()
```

↗

	walk_leisure_past_wk	cigarette_even_once	streets_have_walkways	walk_past_wk	pipe_even_once	walkable_bu
0	1.0	0	0	0	0	
1	1.0	0	1	0	0	
2	1.0	0	1	0	0	
3	1.0	1	1	0	0	
4	0.0	0	1	0	0	

```
df_smoking1['number'] = df_smoking1.index
df_smoking2['number'] = df_smoking2.index

df_smoking1.loc[df_smoking1.number.isin(df_smoking2.number), features1] = df_smoking2[features1]
df_smoking1.head()
```

↗

	language	cereal_serve_per_month	cereal_times_per_month	more_than_one_cereal_type	milk_serve_per_month	
0	5	3	2	0.0	3	
1	4	0	0	0.0	0	
2	5	5	2	0.0	5	
3	3	1	1	0.0	4	
4	5	2	2	1.0	0	

5 rows × 93 columns

```
df_smoking1 = df_smoking1.drop('number', axis = 1)
df_smoking1.head()
```

↗


	language	cereal_serve_per_month	cereal_times_per_month	more_than_one_cereal_type	milk_serve_per_month	
0	5	3	2	0.0	3	
1	4	0	0	0.0	0	
2	5	5	2	0.0	5	
3	3	1	1	0.0	4	
4	5	2	2	1.0	0	

5 rows × 92 columns

```
# Frequeuncy plot for cigarettes_per_day
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline

d = df_smoking1['cigarettes_per_day']
plt.hist(df_smoking1['cigarettes_per_day'], normed=True, bins=15)
plt.ylabel('Frequency');
```

↗ /usr/local/lib/python3.6/dist-packages/matplotlib/axes/\_axes.py:6521: MatplotlibDeprecationWarning:  
The 'normed' kwarg was deprecated in Matplotlib 2.1 and will be removed in 3.1. Use 'density' instead.  
alternative="density", removal="3.1")




```
# Drop rows where cigarettes_per_day = 0
df_smoking1['cigarettes_per_day'] = df_smoking1['cigarettes_per_day'].replace({np.NaN: 0})
df_smoking1 = df_smoking1[df_smoking1['cigarettes_per_day'] > 0]
df_smoking1.shape
```

↳ (7602, 92)

```
# Create frequency plot of cigarettes per day
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline

d = df_smoking1['cigarettes_per_day']
plt.hist(df_smoking1['cigarettes_per_day'], normed=True, bins=15)
plt.ylabel('Frequency');
```

↳ /usr/local/lib/python3.6/dist-packages/matplotlib/axes/\_axes.py:6521: MatplotlibDeprecationWarning:  
The 'normed' kwarg was deprecated in Matplotlib 2.1 and will be removed in 3.1. Use 'density' instead.  
alternative="density", removal="3.1")



```
# Create a column in which cigarettes per day are sorted into 8 bins
df_smoking1['cigarettes_per_day_bins'] = pd.cut(x=df_smoking1['cigarettes_per_day'], bins=[0, 7, 14, 21, 28, 35, 42, 49, 100])
df_smoking1 = df_smoking1.drop('cigarettes_per_day', axis = 1)
df_smoking1['cigarettes_per_day_bins'] = df_smoking1['cigarettes_per_day_bins'].replace ({np.NaN: 0})
df_smoking1.head()
```

↳

	language	cereal_serve_per_month	cereal_times_per_month	more_than_one_cereal_type	milk_serve_per_month
4	5	2	2	1.0	0
9	1	3	2	0.0	1
11	5	0	0	0.0	0
13	5	0	0	0.0	0
14	2	0	0	0.0	0

5 rows × 92 columns

```
# Feature Engineering

# walk_leisure_distance_week = walking_leisure_distance * walk_number_week
df_smoking1['walk_leisure_distance_week'] = df_smoking1['walk_leisure_distance'] * df_smoking1['walk_number_wk']

# single_walk_distance_week = single_walk_distance * walk_number_week
df_smoking1['single_walk_distance_week'] = df_smoking1['single_walk_distance'] * df_smoking1['walk_number_wk']

# tobacco_even_once = cigarette_even_once + cigar_even_once + smokeless_even_once
df_smoking1['tobacco_even_once'] = df_smoking1['cigarette_even_once'] + df_smoking1['cigar_even_once'] + df_smoking1['smokeless_even_once']

# red_meat_eat_serve_per_time = red_meat_eat_serve_month / red_meat_eat_times_month
df_smoking1['red_meat_eat_serve_per_time'] = df_smoking1['red_meat_eat_serve_per_month'] / df_smoking1['red_meat_eat_times_per_month']

# bread_eat_serve_per_time = bread_eat_serve_month / bread_eat_times_month
df_smoking1['bread_eat_serve_per_time'] = df_smoking1['bread_eat_serve_per_month'] / df_smoking1['bread_eat_times_per_month']

df_smoking1.head()
```



↗

	language	cereal_serve_per_month	cereal_times_per_month	more_than_one_cereal_type	milk_serve_per_month
4	5	2	2	1.0	0
9	1	3	2	0.0	1
11	5	0	0	0.0	0
13	5	0	0	0.0	0
14	2	0	0	0.0	0


5 rows × 97 columns

# Looking at the frequency distribution of cigarettes per day bins

```
import matplotlib.pyplot as plt
import numpy as np
%matplotlib inline
```

```
d_bin = df_smoking1['cigarettes_per_day_bins']
plt.hist(d_bin, normed=True, bins=15)
plt.ylabel('Frequency')
```

↗ /usr/local/lib/python3.6/dist-packages/matplotlib/axes/\_axes.py:6521: MatplotlibDeprecationWarning:  
The 'normed' kwarg was deprecated in Matplotlib 2.1 and will be removed in 3.1. Use 'density' instead.  
alternative="density", removal="3.1")  
Text(0, 0.5, 'Frequency')



# Train/validate split: random 80/20% train/validate split.

```
from sklearn.model_selection import train_test_split
XTrain, XVal, yTrain, yVal = train_test_split(df_smoking1.drop('cigarettes_per_day_bins', axis = 1), df_smoking1['cigarettes_per_day_bins'],
                                              test_size=0.2, random_state=42)
XTrain.shape, yTrain.shape, XVal.shape, yVal.shape
```

↗ ((6081, 96), (6081,), (1521, 96), (1521,))

```
# Look at correlation coefficients
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', 1000)
XTrain.corr()
```

↗

	language	cereal_serve_per_month	cereal_times_per_month	more_than_one_cereal_type
language	1.000000	0.436982	0.351576	
cereal_serve_per_month	0.436982	1.000000	0.760684	
cereal_times_per_month	0.351576	0.760684	1.000000	
more_than_one_cereal_type	-0.035361	-0.138573	0.103886	
milk_serve_per_month	0.433675	0.972695	0.735602	
milk_times_per_month	0.349838	0.769347	0.739144	
milk_type	-0.096036	-0.232121	-0.007578	
soda_serve_per_month	0.431958	0.959336	0.721514	
soda_times_per_month	0.342304	0.734191	0.595590	
juice_serve_per_month	0.428804	0.956040	0.720313	
juice_times_per_month	0.332304	0.727421	0.597924	
coffee_serve_per_month	0.426747	0.951691	0.714146	
coffee_times_per_month	0.333119	0.801064	0.622032	
sports_drink_serve_per_month	0.432197	0.957457	0.718121	
sports_drink_times_per_month	0.359200	0.808602	0.625149	
fruit_drink_serve_per_month	0.431355	0.952001	0.713791	
fruit_drink_times_per_month	0.358626	0.798001	0.620712	
fruit_eat_serve_per_month	0.425964	0.957833	0.721305	
fruit_eat_times_per_month	0.384347	0.806646	0.658352	
salad_eat_serve_per_month	0.427673	0.950363	0.713253	
salad_eat_times_per_month	0.382662	0.789765	0.644858	
fries_eat_serve_per_month	0.425416	0.950622	0.710713	
fries_eat_times_per_month	0.361141	0.706499	0.579918	
potatoe_eat_serve_per_month	0.422435	0.936681	0.699211	
potatoe_eat_times_per_month	0.375218	0.743602	0.606211	
beans_eat_serve_per_month	0.421520	0.935026	0.698968	
beans_eat_times_per_month	0.334060	0.704172	0.577761	
grains_eat_serve_per_month	0.422670	0.940141	0.701947	
grains_eat_times_per_month	0.352108	0.698946	0.547232	
vegies_eat_serve_per_month	0.415677	0.928090	0.693861	
vegies_eat_times_per_month	0.359752	0.801514	0.632530	
salsa_eat_serve_per_month	0.421930	0.932706	0.695506	
salsa_eat_times_per_month	0.332938	0.678452	0.541066	
pizza_eat_serve_per_month	0.422585	0.938145	0.699300	
pizza_eat_times_per_month	0.358019	0.679303	0.546140	
tomatoe_eat_serve_per_month	0.418889	0.930008	0.692785	
tomatoe_eat_times_per_month	0.360487	0.700663	0.569326	

cheese_eat_serve_per_month	0.417031	0.926477	0.691735
cheese_eat_times_per_month	0.363737	0.769202	0.610668
red_meat_eat_serve_per_month	0.419657	0.929806	0.694151
red_meat_eat_times_per_month	0.376608	0.780559	0.615793
processed_meat_eat_serve_per_month	0.418972	0.928255	0.692179
processed_meat_eat_times_per_month	0.373554	0.707912	0.571415
bread_eat_serve_per_month	0.417267	0.923150	0.689785
bread_eat_times_per_month	0.339279	0.735331	0.595573
candy_eat_serve_per_month	0.411998	0.922073	0.689743
candy_eat_times_per_month	0.372756	0.707072	0.583550
donut_eat_serve_per_month	0.416284	0.926723	0.690687
donut_eat_times_per_month	0.334741	0.680731	0.556009
cookie_eat_serve_per_month	0.409480	0.912101	0.677290
cookie_eat_times_per_month	0.355908	0.682247	0.559441
ice_cream_eat_serve_per_month	0.414443	0.918537	0.683445
ice_cream_eat_times_per_month	0.350857	0.677407	0.552084
pop_corn_eat_serve_per_month	0.415217	0.921843	0.687277
pop_corn_eat_times_per_month	0.354492	0.669004	0.529327
vitamin_past_month	-0.050629	-0.243404	-0.157238
multivitamin_past_month	-0.037872	-0.162842	-0.096123
multivitamin_days_in_month	-0.029406	-0.150437	-0.089361
calcium_past_month	-0.040267	-0.096730	-0.061498
calcium_days_in_month	-0.034379	-0.086469	-0.060933
vitD_past_month	-0.016192	-0.122617	-0.076643
vitD_days_in_month	-0.013972	-0.111407	-0.068578
vitD_reason	-0.011984	-0.099275	-0.061147
1st_kind_cereal_eaten	-0.066491	-0.213615	0.202229
2nd_kind_cereal_eaten	-0.021112	-0.118378	0.093967
walk_past_wk	-0.100718	-0.114823	-0.085251
walk_number_wk	-0.049873	-0.039521	-0.041604
single_walk_distance	-0.015167	-0.034909	-0.037080
single_walk_time	-0.075258	-0.097345	-0.084728
walk_leisure_past_wk	-0.077325	-0.188538	-0.135776
walk_leisure_number_wk	-0.026543	-0.105001	-0.087298
walk_leisure_distance	-0.026035	-0.067584	-0.044969
walk_leisure_time	-0.061651	-0.163052	-0.120797
see_walking_from_home	0.322965	0.612504	0.441254
weather_discourages_walk	0.214795	0.481079	0.334835
walkway_existence	-0.203418	-0.385381	-0.283120

walkable_retail	-0.159764	-0.199678	-0.134860
walkable_bus_stop	-0.188837	-0.181334	-0.142217
walkable_entertainment	-0.150265	-0.176244	-0.124428
walkable_relaxation	-0.141028	-0.274859	-0.193180
streets_have_walkways	-0.188904	-0.217642	-0.159778
traffic_discourages_walking	-0.093775	-0.097254	-0.076176
crime_discourages_walking	-0.096958	-0.069252	-0.066612
animals_discourage_walking	-0.069518	-0.061819	-0.047632
cigarette_even_once	-0.014661	-0.082766	-0.060123
cigar_even_once	0.017100	-0.156603	-0.099829
pipe_even_once	0.021861	-0.104214	-0.052365
smokeless_even_once	0.036964	-0.087348	-0.057695
had_genetic_counseling	-0.011091	-0.026606	-0.011029
genetic_counseling_with_MD	-0.021622	-0.039074	-0.013490
genetic_counseling_for_cancer	-0.015048	-0.023971	-0.022560
walk_leisure_distance_week	-0.020607	-0.037062	-0.034939
single_walk_distance_week	-0.009937	-0.026258	-0.036861
tobacco_even_once	0.019907	-0.162235	-0.107231
red_meat_eat_serve_per_time	0.428273	0.926809	0.701372
bread_eat_serve_per_time	0.448476	0.925059	0.727886

```
# Dropping highly correlated columns
def correlation(dataset, validation_dataset, threshold):
    col_corr = set() # Set of all the names of deleted columns
    corr_matrix = dataset.corr()
    for i in range(len(corr_matrix.columns)):
        for j in range(i):
            if (corr_matrix.iloc[i, j] >= threshold) and (corr_matrix.columns[j] not in col_corr):
                colname = corr_matrix.columns[i] # getting the name of column
                col_corr.add(colname)
                if colname in dataset.columns:
                    del dataset[colname] # deleting the column from the dataset
                    del validation_dataset[colname] # deleting the column from the validation dataset
```

```
correlation(XTrain, XVal, 0.98)
```

```
XTrain.shape
XVal.shape
```

```
↳ (1521, 81)
```

```
# Begin with baselines for classification.
# The baseline accuracy, if the majority class is guessed for every prediction?
# option with pandas function:
yTrain.value_counts(normalize=True)
```

```
↳
```

```

2      0.286466

# option with scikit-learn function
from sklearn.metrics import accuracy_score
y = yTrain
majority_class = y.mode()[0]
y_pred = [majority_class] * len(y)
accuracy_score(y, y_pred)

0.2864660417694458

# Thus, baseline accuracy, if you guessed the majority class for every prediction is 0.286

# Optimizing Hyperparameters
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

# Define classifier
forest = RandomForestClassifier(random_state = 1)

# Input
X_train = XTrain
y_train = yTrain
X_val = XVal
y_val = yVal

# Parameters to fit
n_estimators = [5, 10, 45, 46, 152, 205, 358, 393, 1000]
max_depth = [3, 5, 7, 10, 15]
min_samples_split = [2, 5, 10]
min_samples_leaf = [1, 5, 10, 15]
max_leaf_nodes = [None, 10, 52]
max_features = [0.11373956383989692, 0.14621091571560108, 0.17046743865886782, 0.17281968473284381, 0.5545636480509806, 0.61]

hyperF = dict(n_estimators = n_estimators, max_depth = max_depth,
              min_samples_split = min_samples_split,
              min_samples_leaf = min_samples_leaf,
              max_leaf_nodes = max_leaf_nodes,
              max_features = max_features)

gridF = GridSearchCV(forest, hyperF, cv = 3, verbose = 10,
                    scoring='accuracy', return_train_score=True,
                    n_jobs = -1)
bestF = gridF.fit(X_train, y_train)

# Output best accuracy and best parameters
print('The score achieved with the best parameters = ', gridF.best_score_, '\n')
print('The parameters are:', gridF.best_params_)

# Use a scikit-learn pipeline to encode categoricals and fit a Random Forest Classifier model.

X_train = XTrain
y_train = yTrain
X_val = XVal
y_val = yVal

from sklearn.pipeline import make_pipeline
import category_encoders as ce
from sklearn.impute import SimpleImputer
from sklearn.ensemble import RandomForestClassifier

pipeline = make_pipeline(
    ce.OneHotEncoder(use_cat_names=True),
    SimpleImputer(strategy='mean'),
    RandomForestClassifier(random_state = 42, max_depth = 10,
                          max_features = 0.11373956383989692,
                          max_leaf_nodes = None,
                          min_samples_leaf = 1,
                          min_samples_split = 10,
                          n_estimators = 205))


pipeline.fit(X_train, y_train)

# Get the model's validation accuracy
ce.OneHotEncoder(use_cat_names=True),
print('Validation Accuracy', pipeline.score(X_val, y_val))

```

Validation Accuracy 0.398422090729783

```
# Plot of features
%matplotlib inline
import matplotlib.pyplot as plt
# Get feature importances
encoder = pipeline.named_steps['onehotencoder']
encoded = encoder.transform(X_train)
rf = pipeline.named_steps['randomforestclassifier']
importances1 = pd.Series(rf.feature_importances_, encoded.columns)
# Plot feature importances
n = 20
plt.figure(figsize=(10,n/2))
plt.title(f'Top {n} features')
importances1.sort_values()[-n:].plot.barh(color='grey');
```




```
# Generate validation curves
%matplotlib inline
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import validation_curve
from sklearn.tree import DecisionTreeClassifier
pipeline = make_pipeline(
    ce.OrdinalEncoder(),
    SimpleImputer(),
    DecisionTreeClassifier()
)

depth = range(1, 10, 2)
train_scores, val_scores = validation_curve(
    pipeline, X_train, y_train,
    param_name='decisiontreeclassifier__max_depth',
    param_range=depth, scoring='accuracy',
    cv=3,
    n_jobs=-1
)
```

```
plt.figure(dpi=150)
plt.plot(depth, np.mean(train_scores, axis=1), color='blue', label='training error')
plt.plot(depth, np.mean(val_scores, axis=1), color='red', label='validation error')
plt.title('Validation Curve')
plt.xlabel('model complexity: RandomForestClassifier max_depth')
plt.ylabel('model score: Accuracy')
plt.legend();
```



## Validation Curve



```
# Tuning the hyper-parameters for a Random Forrest Classifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from scipy.stats import randint, uniform
from sklearn.pipeline import make_pipeline
import category_encoders as ce
from sklearn.impute import SimpleImputer
from sklearn.ensemble import RandomForestClassifier

pipeline = make_pipeline(
    ce.OneHotEncoder(use_cat_names=True),
    SimpleImputer(),
    RandomForestClassifier(random_state = 42, max_depth = 10,
                          max_features = 0.11373956383989692,
                          max_leaf_nodes = None,
                          min_samples_leaf = 1,
                          min_samples_split = 10,
                          n_estimators = 205)
)

param_distributions = {'simpleimputer__strategy': ['mean', 'median', 'most_frequent']}
search = RandomizedSearchCV( pipeline, param_distributions=param_distributions, n_iter=10, cv=3, scoring='accuracy', verbose=
search.fit(X_train, y_train);
```



Fitting 3 folds for each of 3 candidates, totalling 9 fits

/usr/local/lib/python3.6/dist-packages/sklearn/model\_selection/\_search.py:266: UserWarning: The total space of % (grid\_size, self.n\_iter, grid\_size), UserWarning)

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 2 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 1 tasks | elapsed: 3.6s

[Parallel(n\_jobs=-1)]: Done 4 tasks | elapsed: 5.7s

[Parallel(n\_jobs=-1)]: Done 7 out of 9 | elapsed: 9.8s remaining: 2.8s

[Parallel(n\_jobs=-1)]: Done 9 out of 9 | elapsed: 11.2s remaining: 0.0s

[Parallel(n\_jobs=-1)]: Done 9 out of 9 | elapsed: 11.2s finished

```
from sklearn.model_selection import cross_val_score
k = 3
scores = cross_val_score(pipeline, X_val, y_val, cv=k,
scoring='accuracy')
print(f'Validation Accuracy for {k} folds:', scores);
```

➤ Validation Accuracy for 3 folds: [0.39803922 0.39250493 0.39285714]

```
print('Best hyperparameters', search.best_params_)
print('Cross-validation Accuracy', search.best_score_)
```

➤ Best hyperparameters {'simpleimputer\_\_strategy': 'mean'}  
Cross-validation Accuracy 0.3953297155073179


```
pipeline.fit(X_val, y_val)
# Plot of features
%matplotlib inline
import matplotlib.pyplot as plt

# Get feature importances
encoder = pipeline.named_steps['onehotencoder']
encoded = encoder.transform(X_val)
rf = pipeline.named_steps['randomforestclassifier']
importances2 = pd.Series(rf.feature_importances_, encoded.columns)

# Plot feature importances
n = 20
plt.figure(figsize=(10,n/2))
plt.title(f'Top {n} features')
importances2.sort_values()[-n:].plot.barh(color='grey');
```

➤





```
# Demonstrate the relatively high cardinatlity of candy_eat_times_per_month
```

```
XTrain['cookie_eat_serve_per_month'].value_counts()
```



```

1      1730
0      1502
2      1138
3       507
4       265
998     254
5       185
10      120
15       62
7        58
6        57
20       45
8        33
997      32
30       23
999      20
12       17
25       14
18        5
14        4
9         3
203       1
13        1
28        1
24        1
22        1
16        1
31        1

```

Name: cookie\_eat\_serve\_per\_month, dtype: int64

```

# Get drop-column importances
column = 'cookie_eat_serve_per_month'

# # Fit without column
pipeline = make_pipeline(
    ce.OneHotEncoder(use_cat_names=True),
    SimpleImputer(strategy = 'mean'),
    RandomForestClassifier(random_state = 42, max_depth = 10,
                          max_features = 0.11373956383989692,
                          max_leaf_nodes = None,
                          min_samples_leaf = 1,
                          min_samples_split = 10,
                          n_estimators = 205)
)

pipeline.fit(X_train.drop(columns=column), y_train)
score_without = pipeline.score(X_val.drop(columns=column), y_val)
print(f'Validation Accuracy without {column}: {score_without}')

# Fit with column
pipeline = make_pipeline(
    ce.OneHotEncoder(use_cat_names=True),
    SimpleImputer(strategy = 'mean'),
    RandomForestClassifier(random_state = 42, max_depth = 10,
                          max_features = 0.11373956383989692,
                          max_leaf_nodes = None,
                          min_samples_leaf = 1,
                          min_samples_split = 10,
                          n_estimators = 205)
)

pipeline.fit(X_train, y_train)
score_with = pipeline.score(X_val, y_val)
print(f'Validation Accuracy with {column}: {score_with}')

# Compare the error with & without column
print(f'Drop-Column Importance for {column}: {score_with - score_without}')

```

```

↳ Validation Accuracy without cookie_eat_serve_per_month: 0.39316239316239315
Validation Accuracy with cookie_eat_serve_per_month: 0.398422090729783
Drop-Column Importance for cookie_eat_serve_per_month: 0.005259697567389865

```

```
# Rerun the permutation importance process, but for a different feature
```

```

feature = 'language'
X_val_permuted = X_val.copy()
X_val_permuted[feature] = np.random.permutation(X_val[feature])
score_permuted = pipeline.score(X_val_permuted, y_val)

print(f'Validation Accuracy without {feature} permuted: {score_permuted}')
print(f'Validation Accuracy with {feature}: {score_with}')
print(f'Permutation Importance: {score_with - score_permuted}')

```

```

↳ Validation Accuracy without language permuted: 0.3793556870479947
Validation Accuracy with language: 0.398422090729783
Permutation Importance: 0.019066403681788302

```

```

# Using Eli5 library which does not work with pipelines

```

```

transformers = make_pipeline(
    ce.OneHotEncoder(use_cat_names=True),
    SimpleImputer(strategy='mean')
)

```

```

X_train_transformed = transformers.fit_transform(X_train)
X_val_transformed = transformers.transform(X_val)

```

```

model = RandomForestClassifier(random_state = 42, max_depth = 10,
                              max_features = 0.11373956383989692,
                              max_leaf_nodes = None,
                              min_samples_leaf = 1,
                              min_samples_split = 10,
                              n_estimators = 205)

```

```

model.fit(X_train_transformed, y_train)

```

```

↳ RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                          max_depth=10, max_features=0.11373956383989692,
                          max_leaf_nodes=None, min_impurity_decrease=0.0,
                          min_impurity_split=None, min_samples_leaf=1,
                          min_samples_split=10, min_weight_fraction_leaf=0.0,
                          n_estimators=205, n_jobs=None, oob_score=False,
                          random_state=42, verbose=0, warm_start=False)

```

```

# Get permutation importances

```

```

! pip install eli5
from eli5.sklearn import PermutationImportance
import eli5

```

```

permuter = PermutationImportance(
    model,
    scoring='accuracy',
    n_iter=2,
    random_state=42
)

```

```

permuter.fit(X_val_transformed, y_val)
feature_names = X_val.columns.tolist()

```

```

eli5.show_weights(
    permuter,
    top=None, # show permutation importances for all features
    feature_names=feature_names
)

```

```

↳

```

Collecting eli5

Downloading <https://files.pythonhosted.org/packages/97/2f/c85c7d8f8548e460829971785347e14e45fa5c6617da374711c>

|████████████████████| 112kB 10.2MB/s

Requirement already satisfied: jinja2 in /usr/local/lib/python3.6/dist-packages (from eli5) (2.10.3)

Requirement already satisfied: graphviz in /usr/local/lib/python3.6/dist-packages (from eli5) (0.10.1)

Requirement already satisfied: numpy>=1.9.0 in /usr/local/lib/python3.6/dist-packages (from eli5) (1.16.5)

Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from eli5) (1.12.0)

Requirement already satisfied: scikit-learn>=0.18 in /usr/local/lib/python3.6/dist-packages (from eli5) (0.21.3)

Requirement already satisfied: attrs>16.0.0 in /usr/local/lib/python3.6/dist-packages (from eli5) (19.3.0)

Requirement already satisfied: tabulate>=0.7.7 in /usr/local/lib/python3.6/dist-packages (from eli5) (0.8.5)

Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from eli5) (1.3.1)

Requirement already satisfied: MarkupSafe>=0.23 in /usr/local/lib/python3.6/dist-packages (from jinja2->eli5) (1.0.0)

Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.6/dist-packages (from scikit-learn>=0.18->eli5) (0.12.0)

Installing collected packages: eli5

Successfully installed eli5-0.10.1

Using TensorFlow backend.

Weight	Feature
0.0168 ± 0.0033	language
0.0053 ± 0.0039	see_walking_from_home
0.0053 ± 0.0066	bread_eat_serve_per_month
0.0049 ± 0.0007	sports_drink_times_per_month
0.0039 ± 0.0066	walk_leisure_time
0.0036 ± 0.0085	walkable_entertainment
0.0036 ± 0.0007	beans_eat_serve_per_month
0.0033 ± 0.0013	soda_times_per_month
0.0033 ± 0.0013	cookie_eat_times_per_month
0.0030 ± 0.0033	processed_meat_eat_times_per_month
0.0030 ± 0.0007	walkable_bus_stop
0.0030 ± 0.0020	red_meat_eat_times_per_month
0.0026 ± 0.0000	milk_serve_per_month
0.0026 ± 0.0039	tobacco_even_once
0.0026 ± 0.0066	cigar_even_once
0.0026 ± 0.0026	fries_eat_serve_per_month
0.0026 ± 0.0026	salad_eat_times_per_month
0.0023 ± 0.0046	grains_eat_times_per_month
0.0023 ± 0.0033	walk_leisure_distance
0.0023 ± 0.0007	walk_leisure_number_wk
0.0023 ± 0.0072	coffee_times_per_month
0.0020 ± 0.0013	walk_leisure_past_wk
0.0020 ± 0.0013	juice_times_per_month
0.0016 ± 0.0020	walk_past_wk
0.0016 ± 0.0020	vitD_reason
0.0016 ± 0.0007	single_walk_distance
0.0016 ± 0.0007	grains_eat_serve_per_month
0.0016 ± 0.0007	traffic_discourages_walking
0.0013 ± 0.0066	red_meat_eat_serve_per_month
0.0013 ± 0.0013	animals_discourage_walking
0.0013 ± 0.0013	cereal_serve_per_month
0.0013 ± 0.0013	cheese_eat_serve_per_month
0.0010 ± 0.0007	multivitamin_past_month
0.0010 ± 0.0007	cereal_times_per_month
0.0010 ± 0.0007	fries_eat_times_per_month
0.0010 ± 0.0020	ice_cream_eat_times_per_month
0.0010 ± 0.0007	calcium_days_in_month
0.0010 ± 0.0033	walk_number_wk
0.0010 ± 0.0020	vitD_past_month
0.0007 ± 0.0000	walkable_relaxation
0.0007 ± 0.0000	walk_leisure_distance_week
0.0007 ± 0.0013	vitD_days_in_month
0.0007 ± 0.0026	donut_eat_times_per_month
0.0003 ± 0.0007	vitamin_past_month
0.0003 ± 0.0007	genetic_counseling_for_cancer
0.0003 ± 0.0033	multivitamin_days_in_month
0 ± 0.0000	had_genetic_counseling
0 ± 0.0000	genetic_counseling_with_MD
-0.0000 ± 0.0066	pipe_even_once
-0.0000 ± 0.0026	beans_eat_times_per_month
-0.0000 ± 0.0013	more_than_one_cereal_type
-0.0000 ± 0.0039	cookie_eat_serve_per_month
-0.0000 ± 0.0013	tomatoe_eat_times_per_month
-0.0000 ± 0.0007	crime_discourages_walking

```

0.0000 ± 0.0007 smoke_discourages_walking
-0.0003 ± 0.0020 fruit_eat_times_per_month
-0.0007 ± 0.0013 calcium_past_month
-0.0007 ± 0.0013 potatoe_eat_times_per_month
-0.0007 ± 0.0026 weather_discourages_walk
-0.0007 ± 0.0026 candy_eat_times_per_month
-0.0007 ± 0.0026 cheese_eat_times_per_month
-0.0007 ± 0.0013 pop_corn_eat_times_per_month
-0.0010 ± 0.0007 single_walk_distance_week
-0.0010 ± 0.0007 walkway_existence
-0.0010 ± 0.0046 vegies_eat_serve_per_month
-0.0010 ± 0.0007 tomatoe_eat_serve_per_month
-0.0010 ± 0.0007 salsa_eat_times_per_month
-0.0013 ± 0.0026 single_walk_time
-0.0016 ± 0.0033 cigarette_even_once
-0.0016 ± 0.0020 fruit_drink_times_per_month
-0.0023 ± 0.0033 streets_have_walkways
-0.0026 ± 0.0013 2nd_kind_cereal_eaten
-0.0026 ± 0.0053 milk_type
-0.0026 ± 0.0026 pizza_eat_times_per_month
-0.0026 ± 0.0039 vegies_eat_times_per_month
-0.0026 ± 0.0026 bread_eat_times_per_month
-0.0030 ± 0.0085 soda_serve_per_month
-0.0030 ± 0.0020 candy_eat_serve_per_month
-0.0030 ± 0.0046 1st_kind_cereal_eaten
-0.0033 ± 0.0026 milk_times_per_month
-0.0036 ± 0.0020 smokeless_even_once
-0.0036 ± 0.0099 walkable_retail

```

```

# Thus, language is way more important according to feature permutation than according to feature importance in the Random F
# Use importances for feature selection
print('Shape before removing features:', X_train.shape)

```

```

↳ Shape before removing features: (6081, 81)

```

```

# Remove features of 0 importance
zero_importance = 0.0003
mask = permutter.feature_importances_ > zero_importance
features = X_train.columns[mask]
X_train = X_train[features]
print('Shape after removing features:', X_train.shape)

```

```

↳ Shape after removing features: (6081, 46)

```

```

# Random Forest with reduced features to 46
X_val = X_val[features]

pipeline = make_pipeline(
    ce.OneHotEncoder(use_cat_names=True),
    SimpleImputer(strategy='mean'),
    RandomForestClassifier(random_state = 42, max_depth = 10,
                           max_features = 0.11373956383989692,
                           max_leaf_nodes = None,
                           min_samples_leaf = 1,
                           min_samples_split = 10,
                           n_estimators = 205)
)

# Fit on train, score on val
pipeline.fit(X_train, y_train)
print('Validation Accuracy', pipeline.score(X_val, y_val))

```

```

↳ Validation Accuracy 0.4049967126890204

```

```

# Validation Accuracy History
# 0.2864660417694458- baseline guessing the majority class
# 0.4010853478046374- initial fit with optimal hyperparameters

```

```
# 0.398422090729783 - use pipeline with random forest -- with engineered features (0.398422090729783)
# 0.3953297155073179- from cross validation -- with engineered features (0.3953297155073179)
# 0.398422090729783 - doing permutation importance -- with engineered features (0.398422090729783)
# 0.4049967126890204- after removing features of zero importance -- with engineered features (0.4049967126890204)
```

```
# Recursive Feature Elimination
from sklearn.feature_selection import RFECV
from sklearn.model_selection import StratifiedKFold

rfc = RandomForestClassifier(random_state = 42, max_depth = 10,
                           max_features = 0.11373956383989692,
                           max_leaf_nodes = None,
                           min_samples_leaf = 1,
                           min_samples_split = 10,
                           n_estimators = 205)

rfecv = RFECV(estimator=rfc, step=1, cv=StratifiedKFold(9), scoring='accuracy', verbose = 10)
rfecv.fit(X_train, y_train)
```



Fitting estimator with 46 features.  
Fitting estimator with 45 features.  
Fitting estimator with 44 features.  
Fitting estimator with 43 features.  
Fitting estimator with 42 features.  
Fitting estimator with 41 features.  
Fitting estimator with 40 features.  
Fitting estimator with 39 features.  
Fitting estimator with 38 features.  
Fitting estimator with 37 features.  
Fitting estimator with 36 features.  
Fitting estimator with 35 features.  
Fitting estimator with 34 features.  
Fitting estimator with 33 features.  
Fitting estimator with 32 features.  
Fitting estimator with 31 features.  
Fitting estimator with 30 features.  
Fitting estimator with 29 features.  
Fitting estimator with 28 features.  
Fitting estimator with 27 features.  
Fitting estimator with 26 features.  
Fitting estimator with 25 features.  
Fitting estimator with 24 features.  
Fitting estimator with 23 features.  
Fitting estimator with 22 features.  
Fitting estimator with 21 features.  
Fitting estimator with 20 features.  
Fitting estimator with 19 features.  
Fitting estimator with 18 features.  
Fitting estimator with 17 features.  
Fitting estimator with 16 features.  
Fitting estimator with 15 features.  
Fitting estimator with 14 features.  
Fitting estimator with 13 features.  
Fitting estimator with 12 features.  
Fitting estimator with 11 features.  
Fitting estimator with 10 features.  
Fitting estimator with 9 features.  
Fitting estimator with 8 features.  
Fitting estimator with 7 features.  
Fitting estimator with 6 features.  
Fitting estimator with 5 features.  
Fitting estimator with 4 features.  
Fitting estimator with 3 features.  
Fitting estimator with 2 features.  
Fitting estimator with 46 features.  
Fitting estimator with 45 features.  
Fitting estimator with 44 features.  
Fitting estimator with 43 features.  
Fitting estimator with 42 features.  
Fitting estimator with 41 features.  
Fitting estimator with 40 features.  
Fitting estimator with 39 features.  
Fitting estimator with 38 features.  
Fitting estimator with 37 features.  
Fitting estimator with 36 features.  
Fitting estimator with 35 features.  
Fitting estimator with 34 features.  
Fitting estimator with 33 features.  
Fitting estimator with 32 features.  
Fitting estimator with 31 features.  
Fitting estimator with 30 features.  
Fitting estimator with 29 features.  
Fitting estimator with 28 features.  
Fitting estimator with 27 features.  
Fitting estimator with 26 features.  
Fitting estimator with 25 features.  
Fitting estimator with 24 features.  
Fitting estimator with 23 features.  
Fitting estimator with 22 features.  
Fitting estimator with 21 features.

Fitting estimator with 21 features.  
Fitting estimator with 20 features.  
Fitting estimator with 19 features.  
Fitting estimator with 18 features.  
Fitting estimator with 17 features.  
Fitting estimator with 16 features.  
Fitting estimator with 15 features.  
Fitting estimator with 14 features.  
Fitting estimator with 13 features.  
Fitting estimator with 12 features.  
Fitting estimator with 11 features.  
Fitting estimator with 10 features.  
Fitting estimator with 9 features.  
Fitting estimator with 8 features.  
Fitting estimator with 7 features.  
Fitting estimator with 6 features.  
Fitting estimator with 5 features.  
Fitting estimator with 4 features.  
Fitting estimator with 3 features.  
Fitting estimator with 2 features.  
Fitting estimator with 46 features.  
Fitting estimator with 45 features.  
Fitting estimator with 44 features.  
Fitting estimator with 43 features.  
Fitting estimator with 42 features.  
Fitting estimator with 41 features.  
Fitting estimator with 40 features.  
Fitting estimator with 39 features.  
Fitting estimator with 38 features.  
Fitting estimator with 37 features.  
Fitting estimator with 36 features.  
Fitting estimator with 35 features.  
Fitting estimator with 34 features.  
Fitting estimator with 33 features.  
Fitting estimator with 32 features.  
Fitting estimator with 31 features.  
Fitting estimator with 30 features.  
Fitting estimator with 29 features.  
Fitting estimator with 28 features.  
Fitting estimator with 27 features.  
Fitting estimator with 26 features.  
Fitting estimator with 25 features.  
Fitting estimator with 24 features.  
Fitting estimator with 23 features.  
Fitting estimator with 22 features.  
Fitting estimator with 21 features.  
Fitting estimator with 20 features.  
Fitting estimator with 19 features.  
Fitting estimator with 18 features.  
Fitting estimator with 17 features.  
Fitting estimator with 16 features.  
Fitting estimator with 15 features.  
Fitting estimator with 14 features.  
Fitting estimator with 13 features.  
Fitting estimator with 12 features.  
Fitting estimator with 11 features.  
Fitting estimator with 10 features.  
Fitting estimator with 9 features.  
Fitting estimator with 8 features.  
Fitting estimator with 7 features.  
Fitting estimator with 6 features.  
Fitting estimator with 5 features.  
Fitting estimator with 4 features.  
Fitting estimator with 3 features.  
Fitting estimator with 2 features.  
Fitting estimator with 46 features.  
Fitting estimator with 45 features.  
Fitting estimator with 44 features.  
Fitting estimator with 43 features.  
Fitting estimator with 42 features.  
Fitting estimator with 41 features.  
Fitting estimator with 40 features.



Fitting estimator with 39 features.  
Fitting estimator with 38 features.  
Fitting estimator with 37 features.  
Fitting estimator with 36 features.  
Fitting estimator with 35 features.  
Fitting estimator with 34 features.  
Fitting estimator with 33 features.  
Fitting estimator with 32 features.  
Fitting estimator with 31 features.  
Fitting estimator with 30 features.  
Fitting estimator with 29 features.  
Fitting estimator with 28 features.  
Fitting estimator with 27 features.  
Fitting estimator with 26 features.  
Fitting estimator with 25 features.  
Fitting estimator with 24 features.  
Fitting estimator with 23 features.  
Fitting estimator with 22 features.  
Fitting estimator with 21 features.  
Fitting estimator with 20 features.  
Fitting estimator with 19 features.  
Fitting estimator with 18 features.  
Fitting estimator with 17 features.  
Fitting estimator with 16 features.  
Fitting estimator with 15 features.  
Fitting estimator with 14 features.  
Fitting estimator with 13 features.  
Fitting estimator with 12 features.  
Fitting estimator with 11 features.  
Fitting estimator with 10 features.  
Fitting estimator with 9 features.  
Fitting estimator with 8 features.  
Fitting estimator with 7 features.  
Fitting estimator with 6 features.  
Fitting estimator with 5 features.  
Fitting estimator with 4 features.  
Fitting estimator with 3 features.  
Fitting estimator with 2 features.  
Fitting estimator with 46 features.  
Fitting estimator with 45 features.  
Fitting estimator with 44 features.  
Fitting estimator with 43 features.  
Fitting estimator with 42 features.  
Fitting estimator with 41 features.  
Fitting estimator with 40 features.  
Fitting estimator with 39 features.  
Fitting estimator with 38 features.  
Fitting estimator with 37 features.  
Fitting estimator with 36 features.  
Fitting estimator with 35 features.  
Fitting estimator with 34 features.  
Fitting estimator with 33 features.  
Fitting estimator with 32 features.  
Fitting estimator with 31 features.  
Fitting estimator with 30 features.  
Fitting estimator with 29 features.  
Fitting estimator with 28 features.  
Fitting estimator with 27 features.  
Fitting estimator with 26 features.  
Fitting estimator with 25 features.  
Fitting estimator with 24 features.  
Fitting estimator with 23 features.  
Fitting estimator with 22 features.  
Fitting estimator with 21 features.  
Fitting estimator with 20 features.  
Fitting estimator with 19 features.  
Fitting estimator with 18 features.  
Fitting estimator with 17 features.  
Fitting estimator with 16 features.  
Fitting estimator with 15 features.  
Fitting estimator with 14 features.

Fitting estimator with 13 features.  
Fitting estimator with 12 features.  
Fitting estimator with 11 features.  
Fitting estimator with 10 features.  
Fitting estimator with 9 features.  
Fitting estimator with 8 features.  
Fitting estimator with 7 features.  
Fitting estimator with 6 features.  
Fitting estimator with 5 features.  
Fitting estimator with 4 features.  
Fitting estimator with 3 features.  
Fitting estimator with 2 features.  
Fitting estimator with 46 features.  
Fitting estimator with 45 features.  
Fitting estimator with 44 features.  
Fitting estimator with 43 features.  
Fitting estimator with 42 features.  
Fitting estimator with 41 features.  
Fitting estimator with 40 features.  
Fitting estimator with 39 features.  
Fitting estimator with 38 features.  
Fitting estimator with 37 features.  
Fitting estimator with 36 features.  
Fitting estimator with 35 features.  
Fitting estimator with 34 features.  
Fitting estimator with 33 features.  
Fitting estimator with 32 features.  
Fitting estimator with 31 features.  
Fitting estimator with 30 features.  
Fitting estimator with 29 features.  
Fitting estimator with 28 features.  
Fitting estimator with 27 features.  
Fitting estimator with 26 features.  
Fitting estimator with 25 features.  
Fitting estimator with 24 features.  
Fitting estimator with 23 features.  
Fitting estimator with 22 features.  
Fitting estimator with 21 features.  
Fitting estimator with 20 features.  
Fitting estimator with 19 features.  
Fitting estimator with 18 features.  
Fitting estimator with 17 features.  
Fitting estimator with 16 features.  
Fitting estimator with 15 features.  
Fitting estimator with 14 features.  
Fitting estimator with 13 features.  
Fitting estimator with 12 features.  
Fitting estimator with 11 features.  
Fitting estimator with 10 features.  
Fitting estimator with 9 features.  
Fitting estimator with 8 features.  
Fitting estimator with 7 features.  
Fitting estimator with 6 features.  
Fitting estimator with 5 features.  
Fitting estimator with 4 features.  
Fitting estimator with 3 features.  
Fitting estimator with 2 features.  
Fitting estimator with 46 features.  
Fitting estimator with 45 features.  
Fitting estimator with 44 features.  
Fitting estimator with 43 features.  
Fitting estimator with 42 features.  
Fitting estimator with 41 features.  
Fitting estimator with 40 features.  
Fitting estimator with 39 features.  
Fitting estimator with 38 features.  
Fitting estimator with 37 features.  
Fitting estimator with 36 features.  
Fitting estimator with 35 features.  
Fitting estimator with 34 features.  
Fitting estimator with 33 features.

Fitting estimator with 32 features.  
Fitting estimator with 31 features.  
Fitting estimator with 30 features.  
Fitting estimator with 29 features.  
Fitting estimator with 28 features.  
Fitting estimator with 27 features.  
Fitting estimator with 26 features.  
Fitting estimator with 25 features.  
Fitting estimator with 24 features.  
Fitting estimator with 23 features.  
Fitting estimator with 22 features.  
Fitting estimator with 21 features.  
Fitting estimator with 20 features.  
Fitting estimator with 19 features.  
Fitting estimator with 18 features.  
Fitting estimator with 17 features.  
Fitting estimator with 16 features.  
Fitting estimator with 15 features.  
Fitting estimator with 14 features.  
Fitting estimator with 13 features.  
Fitting estimator with 12 features.  
Fitting estimator with 11 features.  
Fitting estimator with 10 features.  
Fitting estimator with 9 features.  
Fitting estimator with 8 features.  
Fitting estimator with 7 features.  
Fitting estimator with 6 features.  
Fitting estimator with 5 features.  
Fitting estimator with 4 features.  
Fitting estimator with 3 features.  
Fitting estimator with 2 features.  
Fitting estimator with 46 features.  
Fitting estimator with 45 features.  
Fitting estimator with 44 features.  
Fitting estimator with 43 features.  
Fitting estimator with 42 features.  
Fitting estimator with 41 features.  
Fitting estimator with 40 features.  
Fitting estimator with 39 features.  
Fitting estimator with 38 features.  
Fitting estimator with 37 features.  
Fitting estimator with 36 features.  
Fitting estimator with 35 features.  
Fitting estimator with 34 features.  
Fitting estimator with 33 features.  
Fitting estimator with 32 features.  
Fitting estimator with 31 features.  
Fitting estimator with 30 features.  
Fitting estimator with 29 features.  
Fitting estimator with 28 features.  
Fitting estimator with 27 features.  
Fitting estimator with 26 features.  
Fitting estimator with 25 features.  
Fitting estimator with 24 features.  
Fitting estimator with 23 features.  
Fitting estimator with 22 features.  
Fitting estimator with 21 features.  
Fitting estimator with 20 features.  
Fitting estimator with 19 features.  
Fitting estimator with 18 features.  
Fitting estimator with 17 features.  
Fitting estimator with 16 features.  
Fitting estimator with 15 features.  
Fitting estimator with 14 features.  
Fitting estimator with 13 features.  
Fitting estimator with 12 features.  
Fitting estimator with 11 features.  
Fitting estimator with 10 features.  
Fitting estimator with 9 features.  
Fitting estimator with 8 features.  
Fitting estimator with 7 features.


```

Fitting estimator with 6 features.
Fitting estimator with 5 features.
Fitting estimator with 4 features.
Fitting estimator with 3 features.
Fitting estimator with 2 features.
Fitting estimator with 46 features.
Fitting estimator with 45 features.
Fitting estimator with 44 features.
Fitting estimator with 43 features.
Fitting estimator with 42 features.
Fitting estimator with 41 features.
Fitting estimator with 40 features.
Fitting estimator with 39 features.
Fitting estimator with 38 features.
Fitting estimator with 37 features.
Fitting estimator with 36 features.
Fitting estimator with 35 features.
Fitting estimator with 34 features.
Fitting estimator with 33 features.
Fitting estimator with 32 features.
Fitting estimator with 31 features.
Fitting estimator with 30 features.
Fitting estimator with 29 features.
Fitting estimator with 28 features.
Fitting estimator with 27 features.
Fitting estimator with 26 features.
Fitting estimator with 25 features.
Fitting estimator with 24 features.
Fitting estimator with 23 features.
Fitting estimator with 22 features.
Fitting estimator with 21 features.
Fitting estimator with 20 features.
Fitting estimator with 19 features.
Fitting estimator with 18 features.
Fitting estimator with 17 features.
Fitting estimator with 16 features.
Fitting estimator with 15 features.
Fitting estimator with 14 features.
Fitting estimator with 13 features.
Fitting estimator with 12 features.
Fitting estimator with 11 features.
Fitting estimator with 10 features.
Fitting estimator with 9 features.
Fitting estimator with 8 features.
Fitting estimator with 7 features.
Fitting estimator with 6 features.
Fitting estimator with 5 features.
Fitting estimator with 4 features.
Fitting estimator with 3 features.
Fitting estimator with 2 features.
Fitting estimator with 46 features.
Fitting estimator with 45 features.
Fitting estimator with 44 features.
Fitting estimator with 43 features.
RFECV(cv=StratifiedKFold(n_splits=9, random_state=None, shuffle=False),
      estimator=RandomForestClassifier(bootstrap=True, class_weight=None,
                                       criterion='gini', max_depth=10,
                                       max_features=0.11373956383989692,
                                       max_leaf_nodes=None,
                                       min_impurity_decrease=0.0,
                                       min_impurity_split=None,
                                       min_samples_leaf=1, min_samples_split=10,
                                       min_weight_fraction_leaf=0.0,
                                       n_estimators=205, n_jobs=None,
                                       oob_score=False, random_state=42,
                                       verbose=0, warm_start=False),
      min_features_to_select=1, n_jobs=None, scoring='accuracy', step=1,
      verbose=10)

```

```
#Plot the results of RFE
plt.figure(figsize=(16, 9))
plt.title('Recursive Feature Elimination with Cross-Validation', fontsize=18, fontweight='bold', pad=20)
plt.xlabel('Number of features selected', fontsize=14, labelpad=20)
plt.ylabel('% Correct Classification', fontsize=14, labelpad=20)
plt.plot(range(1, len(rfecv.grid_scores_) + 1), rfecv.grid_scores_, color='#303F9F', linewidth=3)

plt.show()
```



```
# Print the optimal number of features and accuracy after RFE
print('Optimal number of features: {}'.format(rfecv.n_features_))
```

```
y_pred = rfecv.predict(X_val)
print('Accuracy = ', accuracy_score(y_val, y_pred))
```

↗ Optimal number of features: 42  
Accuracy = 0.41946088099934253

```
# Note that this is a 46.4% improvement over baseline
```

```
# Drop unimportant features
print(np.where(rfecv.support_ == False)[0])
```

```
X_train.drop(X_train.columns[np.where(rfecv.support_ == False)[0]], axis=1, inplace=True)
X_val.drop(X_val.columns[np.where(rfecv.support_ == False)[0]], axis=1, inplace=True)
```

```
X_val.shape
```

```
[26 29 41 43]
```

```
/usr/local/lib/python3.6/dist-packages/pandas/core/frame.py:3940: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-vs-indexing-copy
errors=errors)
(1521, 42)
```

```
X_train.shape
```

```
(6081, 42)
```

```
#Fit to RFECV data set to confirm the best accuracy score
```

```
pipeline0 = make_pipeline(
    ce.OneHotEncoder(use_cat_names=True),
    SimpleImputer(strategy = 'mean'),
    RandomForestClassifier(random_state = 42, max_depth = 10,
                           max_features = 0.11373956383989692,
                           max_leaf_nodes = None,
                           min_samples_leaf = 1,
                           min_samples_split = 10,
                           n_estimators = 205)
)
```

```
# Fit on train, score on val
pipeline0.fit(X_train, y_train)
print('Validation Accuracy', pipeline0.score(X_val, y_val))
```

```
Validation Accuracy 0.41946088099934253
```

```
# Seeing if feature scaling will improve accuracy
from sklearn.preprocessing import MinMaxScaler
```

```
# Get the numbers for the items to be removed from features above
reduced_features = features.delete([26, 29, 41, 43])
```

```
min_max=MinMaxScaler()
# Scaling down both train and test data set
X_train_minmax=min_max.fit_transform(X_train[reduced_features])
X_val_minmax=min_max.fit_transform(X_val[reduced_features])
```

```
#Fit to the scaled data set
```

```
pipeline1 = make_pipeline(
    ce.OneHotEncoder(use_cat_names=True),
    SimpleImputer(strategy = 'mean'),
    RandomForestClassifier(random_state = 42, max_depth = 10,
                           max_features = 0.11373956383989692,
                           max_leaf_nodes = None,
                           min_samples_leaf = 1,
                           min_samples_split = 10,
                           n_estimators = 205)
)
```

```
# Fit on train, score on val
pipeline1.fit(X_train_minmax, y_train)
print('Validation Accuracy', pipeline1.score(X_val_minmax, y_val))
```

```
Validation Accuracy 0.41354372123602895
```

```
# Since scaling does not improve the accuracy score, it is not implemented.
```

```
# Seeing if feature standardization will improve accuracy
from sklearn.preprocessing import scale
```

```
X_train_scale=scale(X_train[reduced_features])
X_val_scale=scale(X_val[reduced_features])
```

#Fit to the standardized data set

```
pipeline2 = make_pipeline(
    ce.OneHotEncoder(use_cat_names=True),
    SimpleImputer(strategy = 'mean'),
    RandomForestClassifier(random_state = 42, max_depth = 10,
                          max_features = 0.11373956383989692,
                          max_leaf_nodes = None,
                          min_samples_leaf = 1,
                          min_samples_split = 10,
                          n_estimators = 205)
)
```

# Fit on train, score on val

```
pipeline2.fit(X_train_scale, y_train)
print('Validation Accuracy', pipeline2.score(X_val_scale, y_val))
```

➞ Validation Accuracy 0.410913872452334

# Since standardizing does not improve the accuracy score, it is not implemented.

# Gradient boosting using XGboost

```
encoder = ce.OrdinalEncoder()
X_train_encoded = encoder.fit_transform(X_train)
X_val_encoded = encoder.transform(X_val)
X_train.shape, X_val.shape, X_train_encoded.shape, X_val_encoded.shape
```

➞ ((6081, 42), (1521, 42), (6081, 42), (1521, 42))

#XGboost with learning\_rate=0.25

```
from xgboost import XGBClassifier
```

```
eval_set = [(X_train_encoded, y_train),
            (X_val_encoded, y_val)]
```

```
model = XGBClassifier(
    random_state = 42,
    max_depth = 10,
    max_features = 0.11373956383989692,
    max_leaf_nodes = None,
    min_samples_leaf = 1,
    min_samples_split = 10,
    n_estimators = 205,
    learning_rate=0.25,
    n_jobs=-1
)
```

```
model.fit(X_train_encoded, y_train, eval_set=eval_set, eval_metric='merror',
          early_stopping_rounds=50)
```

➞

```
[0] validation_0-merror:0.456997 validation_1-merror:0.65812
Multiple eval metrics have been passed: 'validation_1-merror' will be used for early stopping.
```

```
Will train until validation_1-merror hasn't improved in 50 rounds.
```


```
[1] validation_0-merror:0.356685 validation_1-merror:0.648258
[2] validation_0-merror:0.328071 validation_1-merror:0.634451
[3] validation_0-merror:0.281533 validation_1-merror:0.641026
[4] validation_0-merror:0.258675 validation_1-merror:0.637738
[5] validation_0-merror:0.230719 validation_1-merror:0.628534
[6] validation_0-merror:0.21559 validation_1-merror:0.628534
[7] validation_0-merror:0.199967 validation_1-merror:0.631821
[8] validation_0-merror:0.185496 validation_1-merror:0.627876
[9] validation_0-merror:0.173327 validation_1-merror:0.635108
[10] validation_0-merror:0.155731 validation_1-merror:0.629849
[11] validation_0-merror:0.14422 validation_1-merror:0.641026
[12] validation_0-merror:0.13238 validation_1-merror:0.634451
[13] validation_0-merror:0.124979 validation_1-merror:0.643655
[14] validation_0-merror:0.117744 validation_1-merror:0.639053
[15] validation_0-merror:0.111495 validation_1-merror:0.639711
[16] validation_0-merror:0.105081 validation_1-merror:0.639711
[17] validation_0-merror:0.097188 validation_1-merror:0.634451
[18] validation_0-merror:0.08469 validation_1-merror:0.643655
[19] validation_0-merror:0.077125 validation_1-merror:0.644313
[20] validation_0-merror:0.067752 validation_1-merror:0.652202
[21] validation_0-merror:0.05953 validation_1-merror:0.650888
[22] validation_0-merror:0.054925 validation_1-merror:0.646943
[23] validation_0-merror:0.047525 validation_1-merror:0.643655
[24] validation_0-merror:0.044072 validation_1-merror:0.641026
[25] validation_0-merror:0.039632 validation_1-merror:0.642998
[26] validation_0-merror:0.036014 validation_1-merror:0.635766
[27] validation_0-merror:0.032725 validation_1-merror:0.634451
[28] validation_0-merror:0.031409 validation_1-merror:0.638396
[29] validation_0-merror:0.027627 validation_1-merror:0.641026
[30] validation_0-merror:0.025654 validation_1-merror:0.645628
[31] validation_0-merror:0.022365 validation_1-merror:0.639711
[32] validation_0-merror:0.020062 validation_1-merror:0.641683
[33] validation_0-merror:0.01924 validation_1-merror:0.634451
[34] validation_0-merror:0.017267 validation_1-merror:0.635766
[35] validation_0-merror:0.014636 validation_1-merror:0.636423
[36] validation_0-merror:0.014965 validation_1-merror:0.638396
[37] validation_0-merror:0.01332 validation_1-merror:0.633136
[38] validation_0-merror:0.011511 validation_1-merror:0.631821
[39] validation_0-merror:0.009538 validation_1-merror:0.631164
[40] validation_0-merror:0.007893 validation_1-merror:0.637738
[41] validation_0-merror:0.006907 validation_1-merror:0.635108
[42] validation_0-merror:0.006085 validation_1-merror:0.631164
[43] validation_0-merror:0.005756 validation_1-merror:0.633794
[44] validation_0-merror:0.003618 validation_1-merror:0.633794
[45] validation_0-merror:0.00296 validation_1-merror:0.631821
[46] validation_0-merror:0.00296 validation_1-merror:0.631821
[47] validation_0-merror:0.002138 validation_1-merror:0.637081
[48] validation_0-merror:0.001973 validation_1-merror:0.632479
[49] validation_0-merror:0.001809 validation_1-merror:0.632479
[50] validation_0-merror:0.001316 validation_1-merror:0.633794
[51] validation_0-merror:0.001151 validation_1-merror:0.630506
[52] validation_0-merror:0.000987 validation_1-merror:0.633794
[53] validation_0-merror:0.000987 validation_1-merror:0.629849
[54] validation_0-merror:0.000987 validation_1-merror:0.631821
[55] validation_0-merror:0.000987 validation_1-merror:0.633136
[56] validation_0-merror:0.000987 validation_1-merror:0.629849
[57] validation_0-merror:0.000822 validation_1-merror:0.629191
[58] validation_0-merror:0.000822 validation_1-merror:0.629191
Stopping. Best iteration:
[8] validation_0-merror:0.185496 validation_1-merror:0.627876
```

```
XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bynode=1, colsample_bytree=1, gamma=0,
               learning_rate=0.25, max_delta_step=0, max_depth=10,
               max_features=0.11373956383989692, max_leaf_nodes=None,
               min_child_weight=1, min_samples_leaf=1, min_samples_split=10,
               missing=None, n_estimators=205, n_jobs=-1, nthread=None,
```




```
missing=None, n_estimators=200, n_jobs=-1, random_state=None,
objective='multi:softprob', random_state=42, reg_alpha=0,
reg_lambda=1, scale_pos_weight=1, seed=None, silent=None,
subsample=1, verbosity=1)
```

```
# Plot the results
results = model.evals_result()
train_error = results['validation_0']['merror']
val_error = results['validation_1']['merror']
epoch = range(1, len(train_error)+1)
plt.plot(epoch, train_error, label='Train')
plt.plot(epoch, val_error, label='Validation')
plt.ylabel('Classification Error')
plt.xlabel('Model Complexity (n_estimators)')
# plt.ylim((0.5, 0.7)) # Zoom in
plt.legend();
```



```
# Plot log classification error versus model complexity
import numpy as np
results = model.evals_result()
log_train_error = np.log(results['validation_0']['merror'])
log_val_error = np.log(results['validation_1']['merror'])
epoch = range(1, len(train_error)+1)
plt.plot(epoch, log_train_error, label='Train')
plt.plot(epoch, log_val_error, label='Validation')
plt.ylabel('Log Classification Error')
plt.xlabel('Model Complexity (n_estimators)')
# plt.ylim((-0.75, -0.4)) # Zoom in
plt.legend();
```



```
# Note the Classification Error is minimum at n_estimators = 6 in the above
# This is best scene when using the Zoom In scaling
```

```
#Gradient Boosting R^2
from sklearn.metrics import r2_score
from xgboost import XGBRegressor
```

```
gb = make_pipeline(
    ce.OrdinalEncoder(),
    XGBRegressor(n_estimators=46, objective='reg:squarederror', n_jobs=-1)
)

gb.fit(X_train, y_train)
y_pred = gb.predict(X_val)
from sklearn.metrics import r2_score
from xgboost import XGBRegressor
print('Gradient Boosting R^2', r2_score(y_val, y_pred))
```

```
↳ /usr/local/lib/python3.6/dist-packages/xgboost/core.py:587: FutureWarning: Series.base is deprecated and will be removed in a future version.
    if getattr(data, 'base', None) is not None and \
    Gradient Boosting R^2 0.2734730075382791
```

```
# Getting the value distribution for the language feature
df_smoking1['language'].value_counts()
```

```
↳ 5    5713
   4    1031
   8     213
   3     203
   1     169
   2     138
   6     134
   9         1
   Name: language, dtype: int64
```

```
# Define function to vary the language feature while holding all other features constant
import numpy as np
```

```
def vary_language(model, example):
    print('Vary language, hold other features constant', '\n')
    example = example.copy()
    preds = []
    for lang in range(1, 9, 1):
        example['language'] = lang
        pred = model.predict(example)[0]
        print(f'Predicted cigarettes_per_day_bin: {pred:.3f}%')
        print(example.to_string(), '\n')
        preds.append(pred)
    print('Difference between predictions')
    print(np.diff(preds))
```

```
# Vary the language feature while holding all other features constant for the first row
example1 = X_val.iloc[[0]]
vary_language(gb, example1)
```

```
↳
```

Vary language, hold other features constant

Predicted cigarettes\_per\_day\_bin: 3.090%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	ju
31502	1	2	3	3	0	

Predicted cigarettes\_per\_day\_bin: 3.090%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	ju
31502	2	2	3	3	0	

Predicted cigarettes\_per\_day\_bin: 3.099%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	ju
31502	3	2	3	3	0	

Predicted cigarettes\_per\_day\_bin: 3.171%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	ju
31502	4	2	3	3	0	

Predicted cigarettes\_per\_day\_bin: 3.289%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	ju
31502	5	2	3	3	0	

Predicted cigarettes\_per\_day\_bin: 3.289%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	ju
31502	6	2	3	3	0	

Predicted cigarettes\_per\_day\_bin: 3.289%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	ju
31502	7	2	3	3	0	

Predicted cigarettes\_per\_day\_bin: 3.289%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	ju
31502	8	2	3	3	0	

Difference between predictions

```
[0.      0.00901175 0.07167602 0.11833286 0.      0.
 0.      ]
```

```
# Vary the language feature while holding all other features constant for the second row
example2 = X_val.iloc[[2]]
vary_language(gb, example2)
```



Vary language, hold other features constant

Predicted cigarettes\_per\_day\_bin: 2.755%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	juice_per_month
27082	1	1	2	2	0	juice_per_month

Predicted cigarettes\_per\_day\_bin: 2.755%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	juice_per_month
27082	2	1	2	2	0	juice_per_month

Predicted cigarettes\_per\_day\_bin: 2.764%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	juice_per_month
27082	3	1	2	2	0	juice_per_month

Predicted cigarettes\_per\_day\_bin: 2.862%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	juice_per_month
27082	4	1	2	2	0	juice_per_month

Predicted cigarettes\_per\_day\_bin: 2.926%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	juice_per_month
27082	5	1	2	2	0	juice_per_month

Predicted cigarettes\_per\_day\_bin: 2.926%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	juice_per_month
27082	6	1	2	2	0	juice_per_month

Predicted cigarettes\_per\_day\_bin: 2.926%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	juice_per_month
27082	7	1	2	2	0	juice_per_month

Predicted cigarettes\_per\_day\_bin: 2.926%

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	juice_per_month
27082	8	1	2	2	0	juice_per_month

Difference between predictions

```
[0.      0.00901175 0.09764385 0.06409264 0.      0.
 0.      ]
```

# Plot pair dependency of the language feature for rows 1 and 2

%matplotlib inline

import matplotlib.pyplot as plt

examples = pd.concat([example1, example2])

for lang in range(1, 9, 1):

examples['language'] = lang

preds = gb.predict(examples)

for pred in preds:

plt.scatter(lang, pred, color='grey')

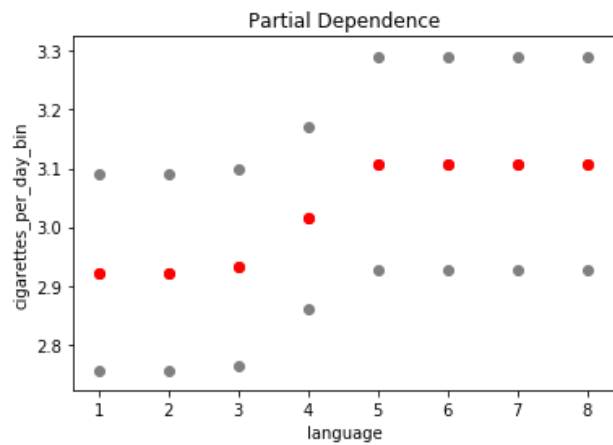
plt.scatter(lang, np.mean(preds), color='red')

plt.title('Partial Dependence')

plt.xlabel('language')

plt.ylabel('cigarettes\_per\_day\_bin')





```
# Create partial dependence plots with one feature
import matplotlib.pyplot as plt
! pip install PDPbox
```

```
# First for the language feature
plt.rcParams['figure.dpi'] = 100
from pdpbox.pdp import pdp_isolate, pdp_plot
feature = 'language'
isolated = pdp_isolate(
    model=gb,
    dataset=X_val,
    model_features=X_val.columns,
    feature=feature
)
pdp_plot(isolated, feature_name=feature);
```




Collecting PDPbox

Downloading <https://files.pythonhosted.org/packages/87/23/ac7da5ba1c6c03a87c412e7e7b6e91a10d6ecf4474906c3e73f6>  
 |██| 57.7MB 1.2MB/s  
 Requirement already satisfied: pandas in /usr/local/lib/python3.6/dist-packages (from PDPbox) (0.24.2)  
 Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from PDPbox) (1.16.5)  
 Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from PDPbox) (1.3.1)  
 Requirement already satisfied: matplotlib>=2.1.2 in /usr/local/lib/python3.6/dist-packages (from PDPbox) (3.0.3)  
 Requirement already satisfied: joblib in /usr/local/lib/python3.6/dist-packages (from PDPbox) (0.14.0)  
 Requirement already satisfied: psutil in /usr/local/lib/python3.6/dist-packages (from PDPbox) (5.4.8)  
 Requirement already satisfied: scikit-learn in /usr/local/lib/python3.6/dist-packages (from PDPbox) (0.21.3)  
 Requirement already satisfied: pytz>=2011k in /usr/local/lib/python3.6/dist-packages (from pandas->PDPbox) (2018.9.2)  
 Requirement already satisfied: python-dateutil>=2.5.0 in /usr/local/lib/python3.6/dist-packages (from pandas->PDPbox) (2.7.3)  
 Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib>=2.1.2->PDPbox) (0.10.0)  
 Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib>=2.1.2->PDPbox) (2.4.6)  
 Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib>=2.1.2->PDPbox) (1.1.0)  
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.6/dist-packages (from python-dateutil>=2.5.0->PDPbox) (1.12.0)  
 Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from kiwisolver>=1.0.1->PDPbox) (44.0.0)  
 Building wheels for collected packages: PDPbox  
 Building wheel for PDPbox (setup.py) ... done  
 Created wheel for PDPbox: filename=PDPbox-0.2.0-cp36-none-any.whl size=57690723 sha256=3a1302daad4c5b733f38f1  
 Stored in directory: /root/.cache/pip/wheels/7d/08/51/63fd122b04a2c87d780464eeffb94867c75bd96a64d500a3fe  
 Successfully built PDPbox  
 Installing collected packages: PDPbox  
 Successfully installed PDPbox-0.2.0

## PDP for feature "language"

Number of unique grid points: 4




```
# Plot partial dependence plot with ICE lines for the language feature
pdp_plot(isolated, feature_name=feature, plot_lines=True, frac_to_plot=100) # Plot 100 ICE lines
plt.xlim(1,8);
```



## PDP for feature "language"

Number of unique grid points: 4

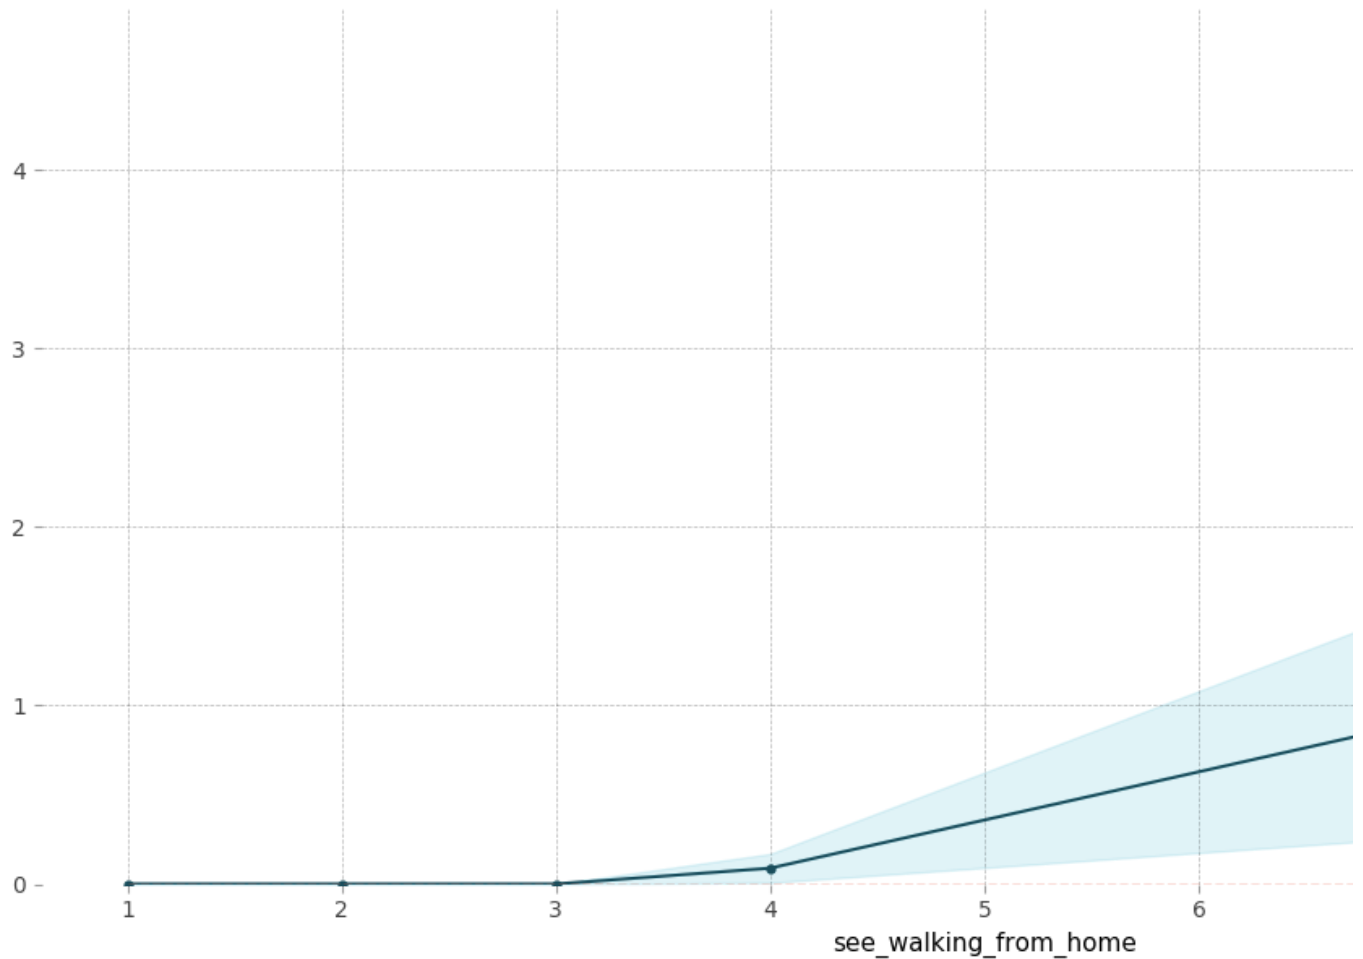


```
# First for the see_walking_from_home feature
plt.rcParams['figure.dpi'] = 100
from pdpbox.pdp import pdp_isolate, pdp_plot
feature = 'see_walking_from_home'
isolated = pdp_isolate(
    model=gb,
    dataset=X_val,
    model_features=X_val.columns,
    feature=feature
)
pdp_plot(isolated, feature_name=feature);
```



## PDP for feature "see\_walking\_from\_home"

Number of unique grid points: 5




```
# Plot partial dependence plot with ICE lines for the see_walking_from_home feature
pdp_plot(isolated, feature_name=feature, plot_lines=True, frac_to_plot=100) # Plot 100 ICE lines
plt.xlim(1,8);
```





## PDP for feature "see\_walking\_from\_home"

Number of unique grid points: 5



```
# Partial Dependence Plots with 2 features
from pdpbox.pdp import pdp_interact, pdp_interact_plot


features = ['language', 'see_walking_from_home']
interaction = pdp_interact(
    model=gb,
    dataset=X_val,
    model_features=X_val.columns,
    features=features
)

pdp_interact_plot(interaction, plot_type='grid', feature_names=features);
```



## PDP interact for "language" and "see\_walking\_from\_home"

Number of unique grid points: (language: 4, see\_walking\_from\_home: 5)




```
# A two feature partial dependence plot in 3D
pdp = interaction.pdp.pivot_table(
    values='preds',
    columns=features[0],
    index=features[1]
)[::-1] # Slice notation to reverse index order so y axis is ascending

import plotly.graph_objs as go

target = 'cigarettes_per_day_bins'

surface = go.Surface(x=pdp.columns,
                    y=pdp.index,
                    z=pdp.values)

layout = go.Layout(
    scene=dict(
        xaxis=dict(title=features[0]),
        yaxis=dict(title=features[1]),
        zaxis=dict(title=target)
    )
)
fig = go.Figure(surface, layout)
fig.show()
```



```
# Test ROC AUC
from sklearn.metrics import roc_auc_score
from sklearn.impute import SimpleImputer
from sklearn.pipeline import make_pipeline
from xgboost import XGBClassifier
import category_encoders as ce

processor = make_pipeline(
    ce.OrdinalEncoder(),
    SimpleImputer(strategy='mean')
)

# Note ROC AUC ranges from 0 - 1, the higher the better
X_val_processed = processor.fit_transform(X_val)

# Contributions to making bin 1 (1 - 7 cigarettes per day) for sample 170
! pip install shap==0.23.0
! pip install -I shap

import shap

row = X_val.iloc[[170]]

explainer = shap.TreeExplainer(model)
row_processed = processor.transform(row)
shap_values_input = explainer.shap_values(row_processed)

shap.initjs()
shap.force_plot(
    base_value=explainer.expected_value[0],
    shap_values=shap_values_input[0],
    features=row
)
```



Collecting shap==0.23.0

Downloading <https://files.pythonhosted.org/packages/60/0d/8bd076821f7230edb2892ad982ea91ca25f2f925466563272ef>  
 |██| 184kB 9.5MB/s  
 Requirement already satisfied: numpy in /usr/local/lib/python3.6/dist-packages (from shap==0.23.0) (1.16.5)  
 Requirement already satisfied: scipy in /usr/local/lib/python3.6/dist-packages (from shap==0.23.0) (1.3.1)  
 Requirement already satisfied: scikit-learn in /usr/local/lib/python3.6/dist-packages (from shap==0.23.0) (0.21)  
 Requirement already satisfied: matplotlib in /usr/local/lib/python3.6/dist-packages (from shap==0.23.0) (3.0.3)  
 Requirement already satisfied: pandas in /usr/local/lib/python3.6/dist-packages (from shap==0.23.0) (0.24.2)  
 Requirement already satisfied: tqdm in /usr/local/lib/python3.6/dist-packages (from shap==0.23.0) (4.28.1)  
 Requirement already satisfied: ipython in /usr/local/lib/python3.6/dist-packages (from shap==0.23.0) (5.5.0)  
 Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.6/dist-packages (from shap==0.23.0) (0.12.0)  
 Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->shap==0.23.0) (1.1.0)  
 Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.6/dist-packages (from matplotlib->shap==0.23.0) (0.10.0)  
 Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->shap==0.23.0) (2.6.1)  
 Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.6/dist-packages (from matplotlib->shap==0.23.0) (2.4.6)  
 Requirement already satisfied: pytz>=2011k in /usr/local/lib/python3.6/dist-packages (from pandas->shap==0.23.0) (2018.9.2)  
 Requirement already satisfied: simplegeneric>=0.8 in /usr/local/lib/python3.6/dist-packages (from ipython->shap==0.23.0) (0.1.2)  
 Requirement already satisfied: prompt-toolkit<2.0.0,>=1.0.4 in /usr/local/lib/python3.6/dist-packages (from ipython->shap==0.23.0) (1.0.15)  
 Requirement already satisfied: pickleshare in /usr/local/lib/python3.6/dist-packages (from ipython->shap==0.23.0) (0.7.5)  
 Requirement already satisfied: decorator in /usr/local/lib/python3.6/dist-packages (from ipython->shap==0.23.0) (4.4.2)  
 Requirement already satisfied: pexpect; sys\_platform != "win32" in /usr/local/lib/python3.6/dist-packages (from ipython->shap==0.23.0) (4.7.0)  
 Requirement already satisfied: setuptools>=18.5 in /usr/local/lib/python3.6/dist-packages (from ipython->shap==0.23.0) (44.0.0)  
 Requirement already satisfied: traitlets>=4.2 in /usr/local/lib/python3.6/dist-packages (from ipython->shap==0.23.0) (4.2.0)  
 Requirement already satisfied: pygments in /usr/local/lib/python3.6/dist-packages (from ipython->shap==0.23.0) (2.2.0)  
 Requirement already satisfied: six in /usr/local/lib/python3.6/dist-packages (from cycler->shap==0.23.0) (1.11.0)  
 Requirement already satisfied: wcwidth in /usr/local/lib/python3.6/dist-packages (from prompt-toolkit<2.0.0,>=1.0.4->shap==0.23.0) (0.1.7)  
 Requirement already satisfied: ptyprocess>=0.5 in /usr/local/lib/python3.6/dist-packages (from pexpect; sys\_platform != "win32"->shap==0.23.0) (0.6.0)  
 Requirement already satisfied: ipython-genutils in /usr/local/lib/python3.6/dist-packages (from traitlets>=4.2->shap==0.23.0) (0.2.0)  
 Building wheels for collected packages: shap  
 Building wheel for shap (setup.py) ... done  
 Created wheel for shap: filename=shap-0.23.0-cp36-cp36m-linux\_x86\_64.whl size=235685 sha256=18c38da919862d09f  
 Stored in directory: /root/.cache/pip/wheels/c1/2c/aa/10d1782fe066536fcd564a2f8adea4dd05f57768236038855b

Successfully built shap

Installing collected packages: shap

Successfully installed shap-0.23.0

Collecting shap

Downloading <https://files.pythonhosted.org/packages/2b/4b/5944c379c94f8f6335dd36b9316292236e3da0dee8da806f60e>  
 |██| 266kB 9.0MB/s

Collecting numpy (from shap)

Downloading <https://files.pythonhosted.org/packages/0e/46/ae6773894f7eac53308086287897ec568eac9768918d913d5f>  
 |██| 20.0MB 50.0MB/s

Collecting scipy (from shap)

Downloading <https://files.pythonhosted.org/packages/29/50/a552a5aff252ae915f522e44642bb49a7b7b31677f9580cfd11>  
 |██| 25.2MB 1.3MB/s

Collecting scikit-learn (from shap)

Downloading <https://files.pythonhosted.org/packages/a0/c5/d2238762d780dde84a20b8c761f563fe882b88c5a5fb03c056f>  
 |██| 6.7MB 43.4MB/s

Collecting pandas (from shap)

Downloading <https://files.pythonhosted.org/packages/86/12/08b092f6fc9e4c2552e37add0861d0e0e0d743f78f1318973c>  
 |██| 10.4MB 34.9MB/s

Collecting tqdm>4.25.0 (from shap)

Downloading <https://files.pythonhosted.org/packages/e1/c1/bc1dba38b48f4ae3c4428aea669c5e27bd5a7642a74c8348451>  
 |██| 61kB 25.4MB/s

Collecting joblib>=0.11 (from scikit-learn->shap)

Downloading <https://files.pythonhosted.org/packages/8f/42/155696f85f344c066e17af287359c9786b436b1bf86029bb341>  
 |██| 296kB 42.6MB/s

Collecting pytz>=2017.2 (from pandas->shap)

Downloading <https://files.pythonhosted.org/packages/e7/f9/f0b53f88060247251bf481fa6ea62cd0d25bf1b11a87888e53c>  
 |██| 512kB 56.8MB/s

Collecting python-dateutil>=2.6.1 (from pandas->shap)

Downloading <https://files.pythonhosted.org/packages/41/17/c62facbfbfd163c7f57f3844689e3a78bae1f403648a6afb1df>  
 |██| 235kB 48.0MB/s

Collecting six>=1.5 (from python-dateutil>=2.6.1->pandas->shap)


Downloading <https://files.pythonhosted.org/packages/73/fb/00a976f728d0d1fecfe898238ce23f502a721c0ac0ecfedb80e>  
 |██| 10.4kB 48.0MB/s

Building wheels for collected packages: shap

Building wheel for shap (setup.py) ... done  
 Created wheel for shap: filename=shap-0.31.0-cp36-cp36m-linux\_x86\_64.whl size=375005 sha256=530f855c4f72a4b5e  
 Stored in directory: /root/.cache/pip/wheels/7b/2d/46/ff8959add2e4e99a18a6e90b82f47508bf52fdf7e7d806f7df  
 Successfully built shap

ERROR: google-colab 1.0.0 has requirement pandas~0.24.0, but you'll have pandas 0.25.2 which is incompatible.  
 ERROR: datascience 0.10.6 has requirement folium==0.2.1, but you'll have folium 0.8.3 which is incompatible.

ERROR: albumentations 0.1.12 has requirement imgaug<0.2.7,>=0.2.5, but you'll have imgaug 0.2.9 which is incompatible.  
 ERROR: albumentations 0.1.12 has requirement imgaug<0.2.7,>=0.2.5, but you'll have imgaug 0.2.9 which is incompatible.  
 Installing collected packages: numpy, scipy, joblib, scikit-learn, pytz, six, python-dateutil, pandas, tqdm, shap  
 Successfully installed joblib-0.14.0 numpy-1.17.3 pandas-0.25.2 python-dateutil-2.8.0 pytz-2019.3 scikit-learn-




```
# Contributions to making bin 8 (49 - more cigarettes per day) for sample 170
import shap

row = X_val.iloc[[170]]

explainer = shap.TreeExplainer(model)
row_processed = processor.transform(row)
shap_values_input = explainer.shap_values(row_processed)

shap.initjs()
shap.force_plot(
    base_value=explainer.expected_value[7],
    shap_values=shap_values_input[7],
    features=row
)
```



```
# Features importances for sample 170

feature_names = row.columns
feature_values = row.values[0]
shap_values_array = np.asarray(shap_values_input)
shaps = pd.Series(shap_values_array[0,0,:], zip(feature_names, feature_values))
shaps.sort_values().plot.barh(color='grey', figsize=(10,15));
```





```
# Create a dataframe for sample 170
# bin versus feature
```

```
my_python_list = [shap_values_array[0, 0, :], shap_values_array[1, 0, :], shap_values_array[2, 0, :], shap_values_array[3, 0, :]]
df_bins = pd.DataFrame(columns=np.array(feature_names), data=my_python_list)

df_bins.head(8)
```

```
↳
```

	language	cereal_serve_per_month	cereal_times_per_month	milk_serve_per_month	soda_times_per_month	juice
0	-0.063460	-0.058806	0.006714	-0.496689	0.086444	
1	-0.017955	-0.160864	0.016525	0.224091	-0.046407	
2	0.087442	0.047602	-0.038343	0.036599	-0.008106	
3	-0.011019	0.011034	-0.000203	0.059331	-0.002250	
4	-0.001570	-0.071335	-0.009197	-0.127074	0.025907	
5	0.005110	0.059201	-0.000520	-0.101339	0.001079	
6	0.000000	0.000542	0.000000	0.000000	0.000000	


```
# Create a 3D plot of force as a function of cigarettes_per_day_bin and feature for sample 170
# A two feature partial dependence plot in 3D
import plotly.graph_objs as go
```

```
surface = go.Surface(x=df_bins.columns,
                    y=df_bins.index + 1,
                    z=df_bins.values)

layout = go.Layout(
    scene=dict(
        xaxis=dict(title= 'Features'),
        yaxis=dict(title= 'cigarettes_per_day_bin'),
        zaxis=dict(title= 'Force')
    )
)

fig = go.Figure(surface, layout)
fig.show()
```

```
↳
```



```

pros = shaps.sort_values(ascending=False)[:3].index
cons = shaps.sort_values(ascending=True)[:3].index

print('Pros:')
for i, pro in enumerate(pros, start=1):
    feature_name, feature_value = pro
    print(f'{i}. {feature_name} is {feature_value}')
print('\n')

print('Cons:')
for i, con in enumerate(cons, start=1):
    feature_name, feature_value = con
    print(f'{i}. {feature_name} is {feature_value}')

```

☞ Pros:

1. soda\_times\_per\_month is 3.0
2. fries\_eat\_times\_per\_month is 3.0
3. cigar\_even\_once is 0.0

Cons:

1. milk\_serve\_per\_month is 25.0
2. walkable\_bus\_stop is 0.0
3. cookie\_eat\_times\_per\_month is 2.0


```

# Create function for constructing confusion matrix
%matplotlib inline
import seaborn as sns
from sklearn.metrics import confusion_matrix
from sklearn.utils.multiclass import unique_labels
def plot_confusion_matrix(y_true, y_pred):
    labels = unique_labels(y_true)
    columns = [f'Predicted {label}' for label in labels]
    index = [f'Actual {label}' for label in labels]
    table = pd.DataFrame(confusion_matrix(y_true, y_pred),
        columns=columns, index=index)
    return sns.heatmap(table, annot=True, fmt='d', cmap='viridis')

```



```
y_pred = pipeline0.predict(X_val)
plot_confusion_matrix(y_val, y_pred);
```



```
# Get precision & recall for majority class baseline
from sklearn.metrics import classification_report
print(classification_report(y_val, y_pred))
```

	precision	recall	f1-score	support
1	0.45	0.63	0.52	452
2	0.25	0.00	0.01	315
3	0.35	0.67	0.46	418
4	0.00	0.00	0.00	17
5	0.00	0.00	0.00	62
6	0.00	0.00	0.00	76
8	0.88	0.41	0.56	181
accuracy			0.42	1521
macro avg	0.28	0.24	0.22	1521
weighted avg	0.39	0.42	0.35	1521

```
# Another way to get a classification report using an ROC_AUC approach (https://stackoverflow.com/questions/39685740/calculating-roc-auc-for-multiclass-problem)
import pandas as pd
import numpy as np
from scipy import interp

from sklearn.metrics import precision_recall_fscore_support
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import LabelBinarizer

def class_report(y_true, y_pred, y_score=None, average='micro'):
    if y_true.shape != y_pred.shape:
        print("Error! y_true %s is not the same shape as y_pred %s" % (
            y_true.shape,
            y_pred.shape))
    )
    return

lb = LabelBinarizer()

if len(y_true.shape) == 1:
    lb.fit(y_true)

#Value counts of predictions
labels, cnt = np.unique(
    y_pred,
    return_counts=True)
n_classes = len(labels)
pred_cnt = pd.Series(cnt, index=labels)

metrics_summary = precision_recall_fscore_support(
    y_true=y_true,
    y_pred=y_pred,
    labels=labels)
```

```

avg = list(precision_recall_fscore_support(
    y_true=y_true,
    y_pred=y_pred,
    average='weighted'))

metrics_sum_index = ['precision', 'recall', 'f1-score', 'support']
class_report_df = pd.DataFrame(
    list(metrics_summary),
    index=metrics_sum_index,
    columns=labels)

support = class_report_df.loc['support']
total = support.sum()
class_report_df['avg / total'] = avg[:-1] + [total]

class_report_df = class_report_df.T
class_report_df['pred'] = pred_cnt
class_report_df['pred'].iloc[-1] = total

if not (y_score is None):
    fpr = dict()
    tpr = dict()
    roc_auc = dict()
    for label_it, label in enumerate(labels):
        fpr[label], tpr[label], _ = roc_curve(
            (y_true == label).astype(int),
            y_score[:, label_it])

        roc_auc[label] = auc(fpr[label], tpr[label])

    if average == 'micro':
        if n_classes <= 2:
            fpr["avg / total"], tpr["avg / total"], _ = roc_curve(
                lb.transform(y_true).ravel(),
                y_score[:, 1].ravel())
        else:
            fpr["avg / total"], tpr["avg / total"], _ = roc_curve(
                lb.transform(y_true).ravel(),
                y_score.ravel())

        roc_auc["avg / total"] = auc(
            fpr["avg / total"],
            tpr["avg / total"])

    elif average == 'macro':
        # First aggregate all false positive rates
        all_fpr = np.unique(np.concatenate([
            fpr[i] for i in labels
        ]))

        # Then interpolate all ROC curves at this points
        mean_tpr = np.zeros_like(all_fpr)
        for i in labels:
            mean_tpr += interp(all_fpr, fpr[i], tpr[i])

        # Finally average it and compute AUC
        mean_tpr /= n_classes

        fpr["macro"] = all_fpr
        tpr["macro"] = mean_tpr

        roc_auc["avg / total"] = auc(fpr["macro"], tpr["macro"])

    class_report_df['AUC'] = pd.Series(roc_auc)

return class_report_df

```

```

# The above function provides the predicted values for each class.
class_report(y_val, y_pred, y_score=None, average='micro')

```



```

precision    recall  f1 score   support
# Deriving an ROC curve for each class in cigarettes_per_day_bins
# Transform y_val and y_pred to arrays that are 1521 by 8 with bins as the columns

y_val_trans = pd.DataFrame(columns=['1', '2', '3', '4', '5', '6', '7', '8'])
y_val_trans['1'] = y_val.map(lambda x : 1 if x==1 else 0)
y_val_trans['2'] = y_val.map(lambda x : 1 if x==2 else 0)
y_val_trans['3'] = y_val.map(lambda x : 1 if x==3 else 0)
y_val_trans['4'] = y_val.map(lambda x : 1 if x==4 else 0)
y_val_trans['5'] = y_val.map(lambda x : 1 if x==5 else 0)
y_val_trans['6'] = y_val.map(lambda x : 1 if x==6 else 0)
y_val_trans['7'] = y_val.map(lambda x : 1 if x==7 else 0)
y_val_trans['8'] = y_val.map(lambda x : 1 if x==8 else 0)
print ('y_val_trans =')
print (y_val_trans.head(), '\n')

y_pred_proba = model.predict_proba(X_val)
y_pred_trans = pd.DataFrame(y_pred_proba)

print ('y_pred_trans')
print (y_pred_trans.head(), '\n')

```

```

In [ ]: y_val_trans =
      1  2  3  4  5  6  7  8
31502  0  0  1  0  0  0  0  0
4439   1  0  0  0  0  0  0  0
27082  0  1  0  0  0  0  0  0
19317  0  1  0  0  0  0  0  0
2063   0  0  0  0  1  0  0  0

y_pred_trans
      0      1      2      3      4      5      6  \
0  0.079808  0.217352  0.328526  0.032976  0.167012  0.070402  0.032201
1  0.189750  0.215191  0.335306  0.048586  0.057663  0.058306  0.046568
2  0.159886  0.216339  0.327539  0.039561  0.086386  0.050499  0.037919
3  0.030064  0.086333  0.076470  0.028838  0.031843  0.028758  0.028475
4  0.227320  0.196754  0.311454  0.039435  0.056393  0.056143  0.037474

      7
0  0.071723
1  0.048630
2  0.081869
3  0.689219
4  0.075027

```

```

# Learn to predict each class against the other
print(__doc__)

import numpy as np

from sklearn import svm, datasets
from sklearn.metrics import roc_curve, auc

# Compute ROC curve and ROC area for each class
fpr = dict()
tpr = dict()
roc_auc = dict()
for i in range(8):
    fpr[i], tpr[i], _ = roc_curve(y_val_trans.iloc[:, i], y_pred_trans.iloc[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

# Compute micro-average ROC curve and ROC area
fpr["micro"], tpr["micro"], _ = roc_curve(y_val_trans.values.ravel(), y_pred_trans.values.ravel())
roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])

```

In [ ]: Automatically created module for IPython interactive environment

```

# Compute macro-average ROC curve and ROC area
import matplotlib.pyplot as plt
from itertools import cycle
from scipy import interp
n_classes = 8

```

```

lw = 2

# First aggregate all false positive rates
all_fpr = np.unique(np.concatenate([fpr[i] for i in range(n_classes)]))

# Then interpolate all ROC curves at this points
mean_tpr = np.zeros_like(all_fpr)
for i in range(n_classes):
    mean_tpr += interp(all_fpr, fpr[i], tpr[i])

# Finally average it and compute AUC
mean_tpr /= n_classes

fpr["macro"] = all_fpr
tpr["macro"] = mean_tpr
roc_auc["macro"] = auc(fpr["macro"], tpr["macro"])

# Plot all ROC curves
plt.figure()
plt.plot(fpr["micro"], tpr["micro"],
         label='micro-average ROC curve (area = {0:0.2f})'
         ''.format(roc_auc["micro"]),
         color='deeppink', linestyle=':', linewidth=4)

plt.plot(fpr["macro"], tpr["macro"],
         label='macro-average ROC curve (area = {0:0.2f})'
         ''.format(roc_auc["macro"]),
         color='navy', linestyle=':', linewidth=4)

colors = cycle(['aqua', 'darkorange', 'cornflowerblue', 'blue', 'green'])
for i, color in zip(range(n_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=lw,
             label='ROC curve of class {0} (area = {1:0.2f})'
             ''.format(i + 1, roc_auc[i]))

plt.plot([0, 1], [0, 1], 'k--', lw=lw)
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Some extension of Receiver operating characteristic to multi-class')
plt.legend(loc="lower right")
plt.show()

```

☞ Some extension of Receiver operating characteristic to multi-class

